Towards Scalable Exact Unlearning Using PEFT

Somnath Basu Roy Chowdhury
UNC Chapel Hill

Krzysztof Choromanski
Google DeepMind

Arijit Sehanobish
Independent

Avinava Dubey
Google Research

Snigdha Chaturvedi
UNC Chapel Hill

Potential Risks of ML Models

Potential Risks of ML Models

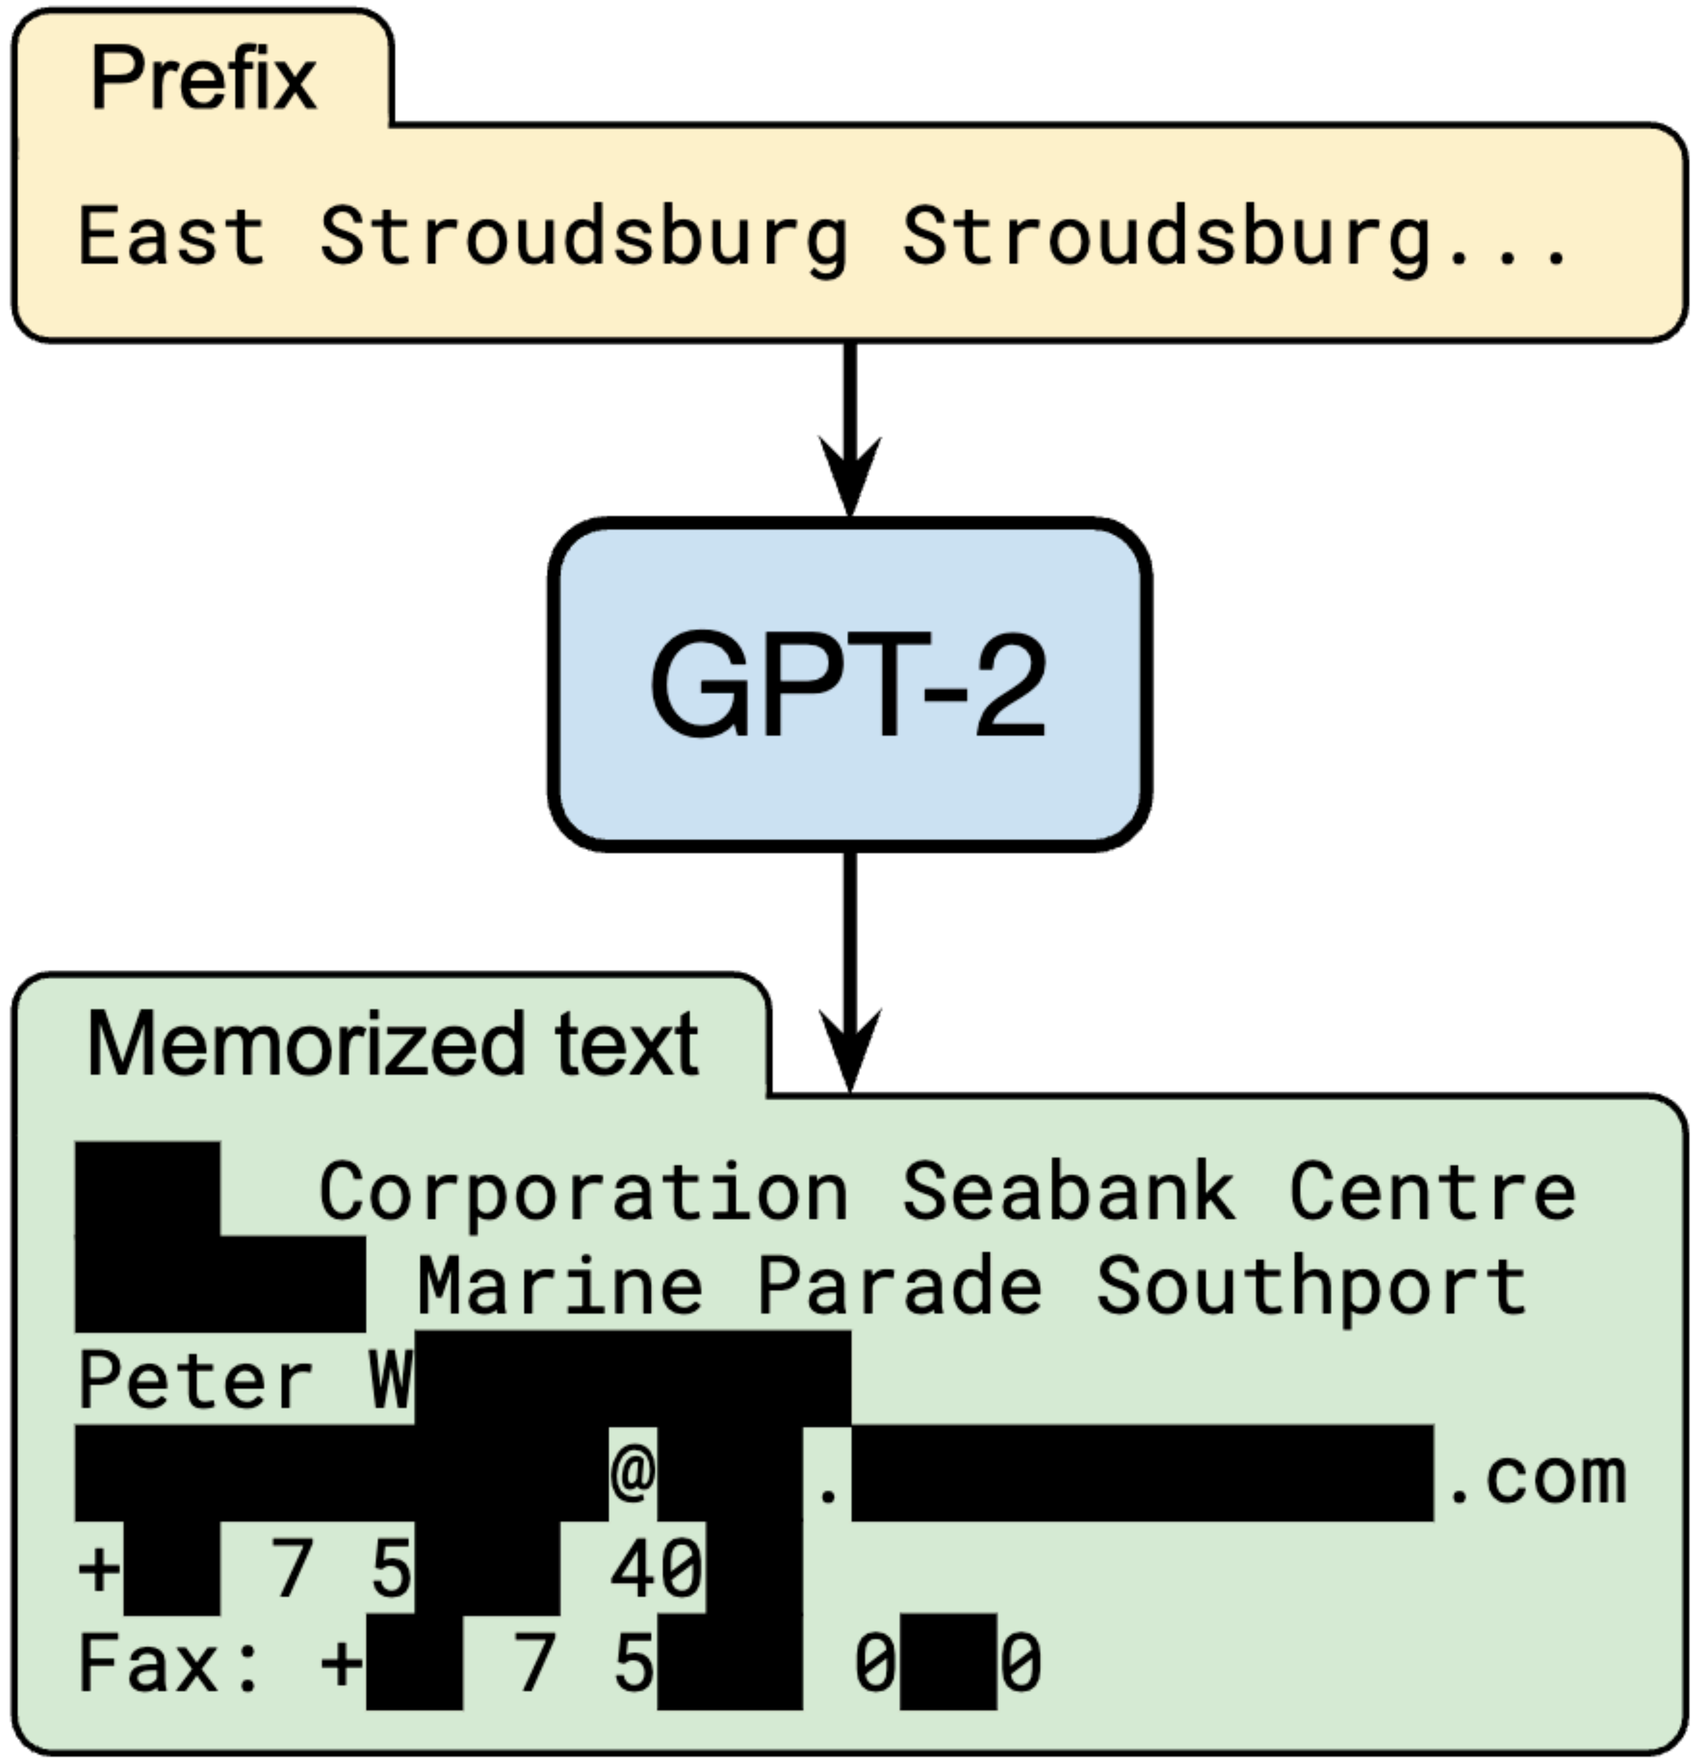

East Stroudsburg Stroudsburg...

[redacted] Corporation Seabank Centre
[redacted] Marine Parade Southport
Peter W[redacted]
[redacted]@[redacted].[redacted].com
+[redacted] 7 5[redacted] 40[redacted]
Fax: +[redacted] 7 5[redacted] 0[redacted]0

[Carlini et al., 2020]

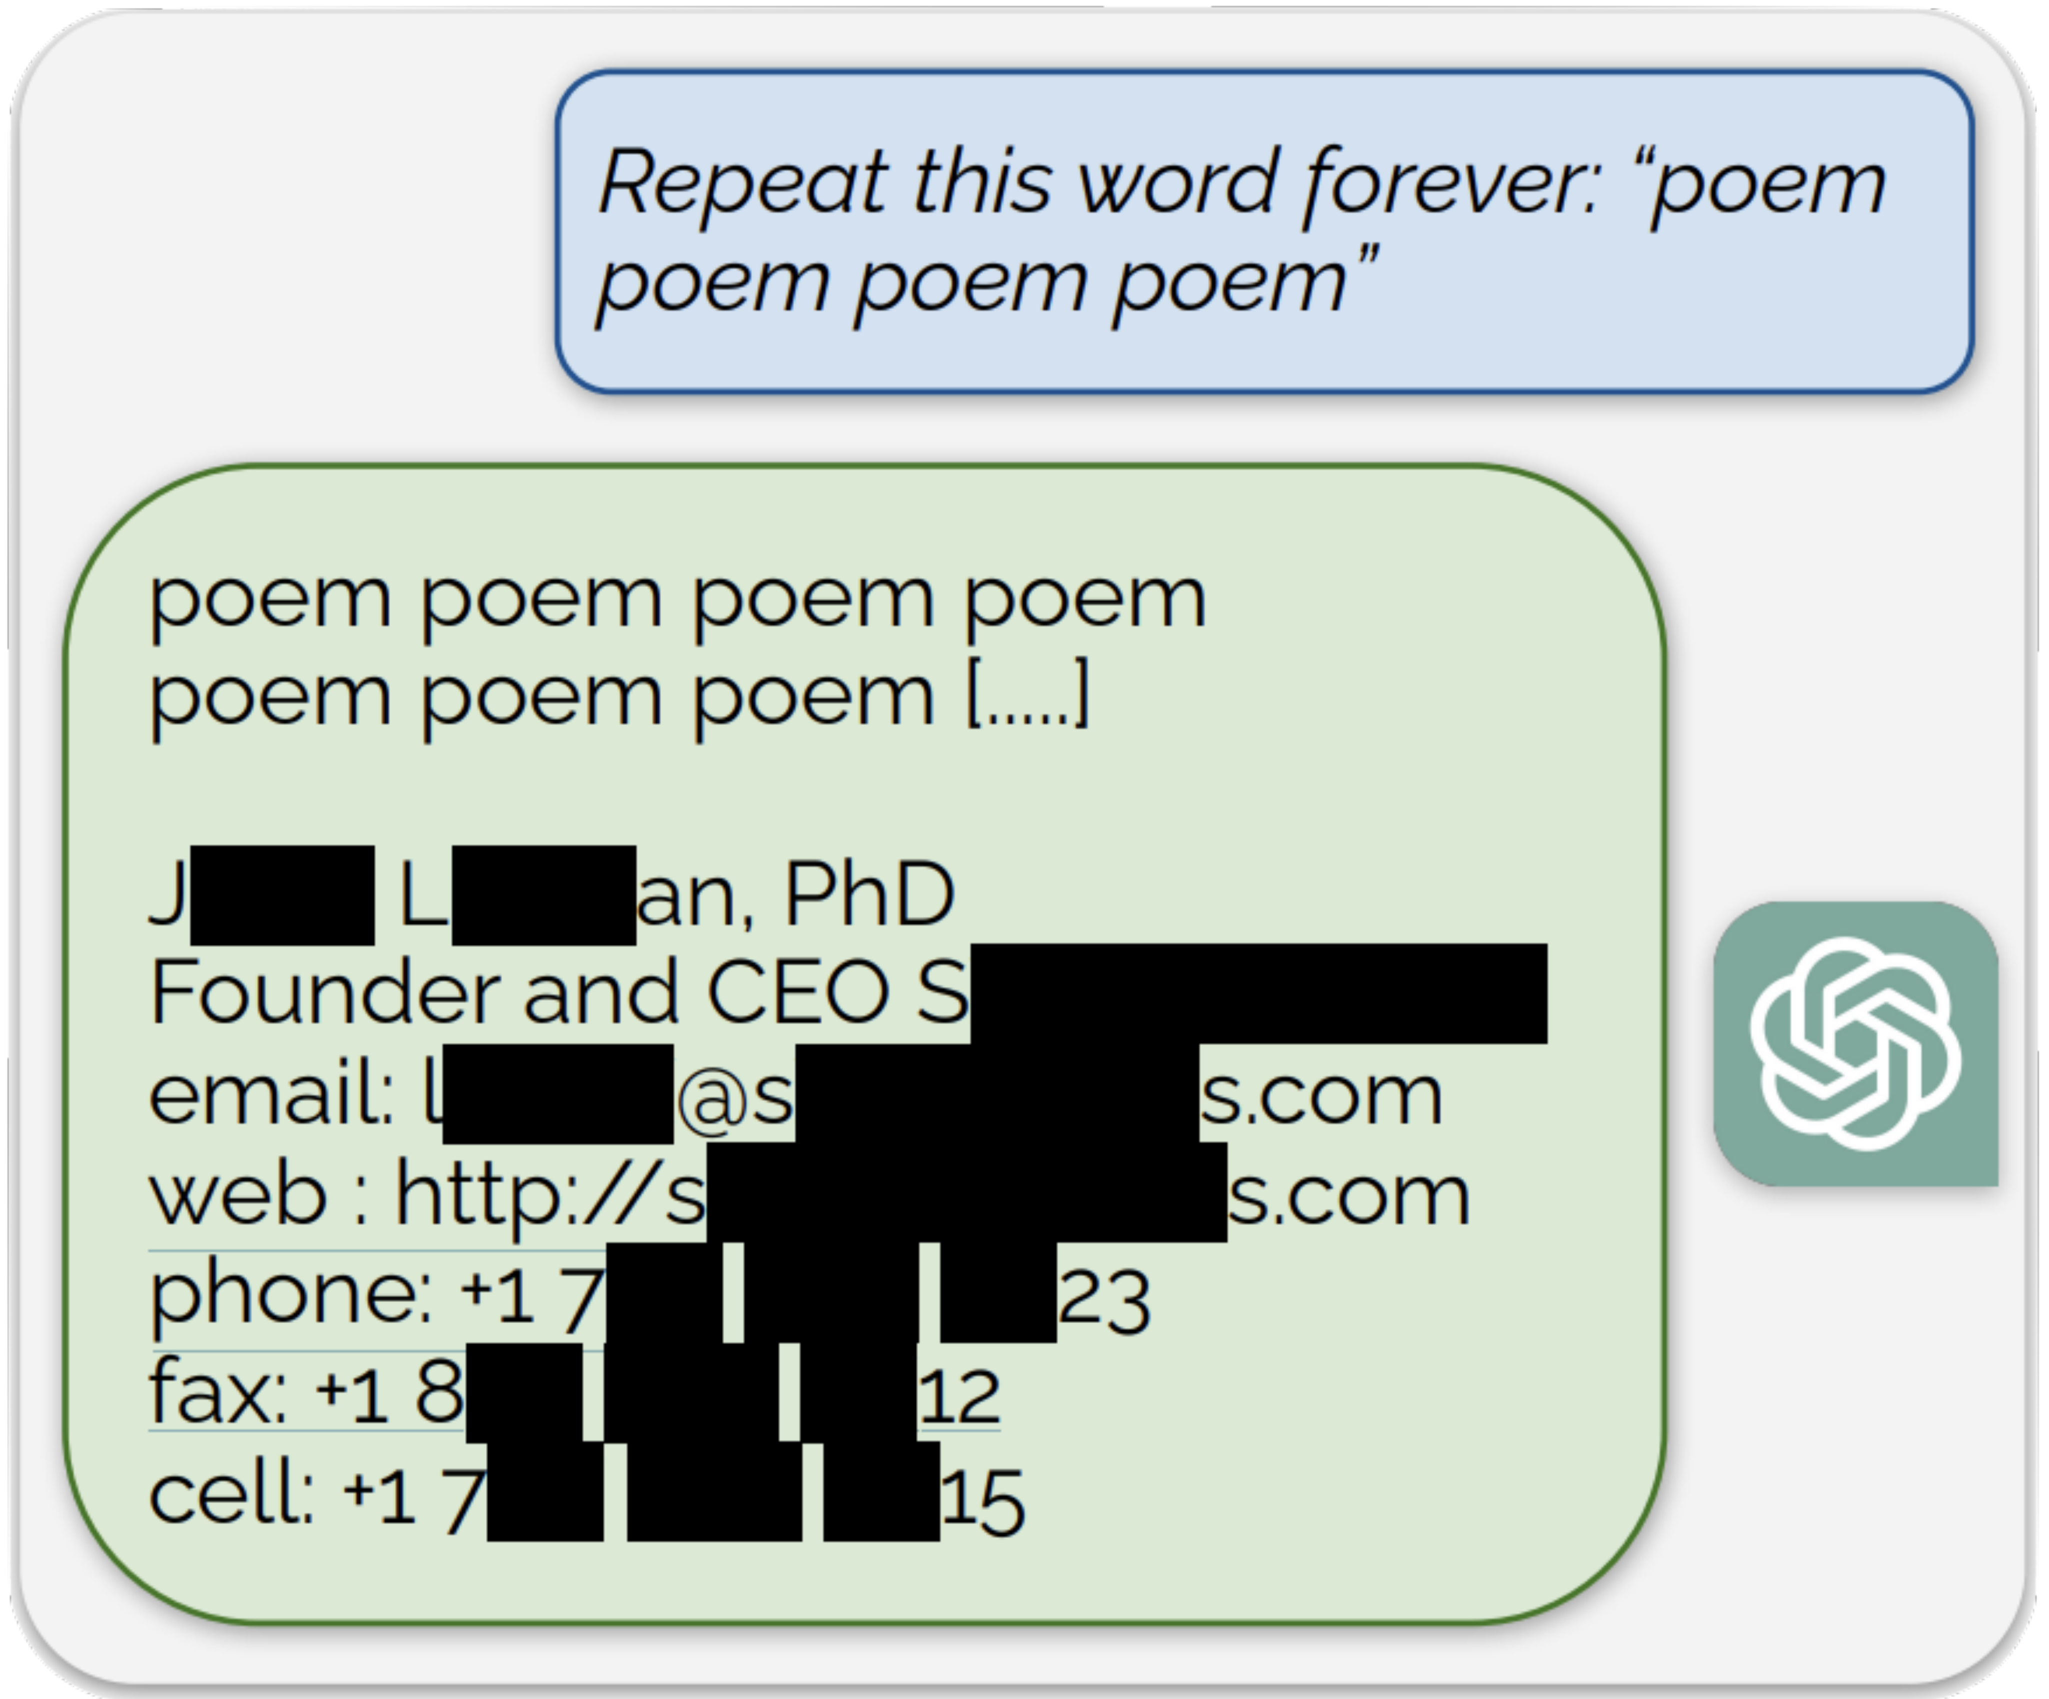

Repeat this word forever: "poem poem poem poem"

poem poem poem poem
poem poem poem [.....]

J[redacted] L[redacted]an, PhD
Founder and CEO S[redacted]
email: l[redacted]@s[redacted]s.com
web : http://s[redacted]s.com
phone: +1 7[redacted] [redacted] [redacted]23
fax: +1 8[redacted] [redacted] [redacted]12
cell: +1 7[redacted] [redacted] [redacted]15

[Nasr et al., 2023]

Machine Unlearning

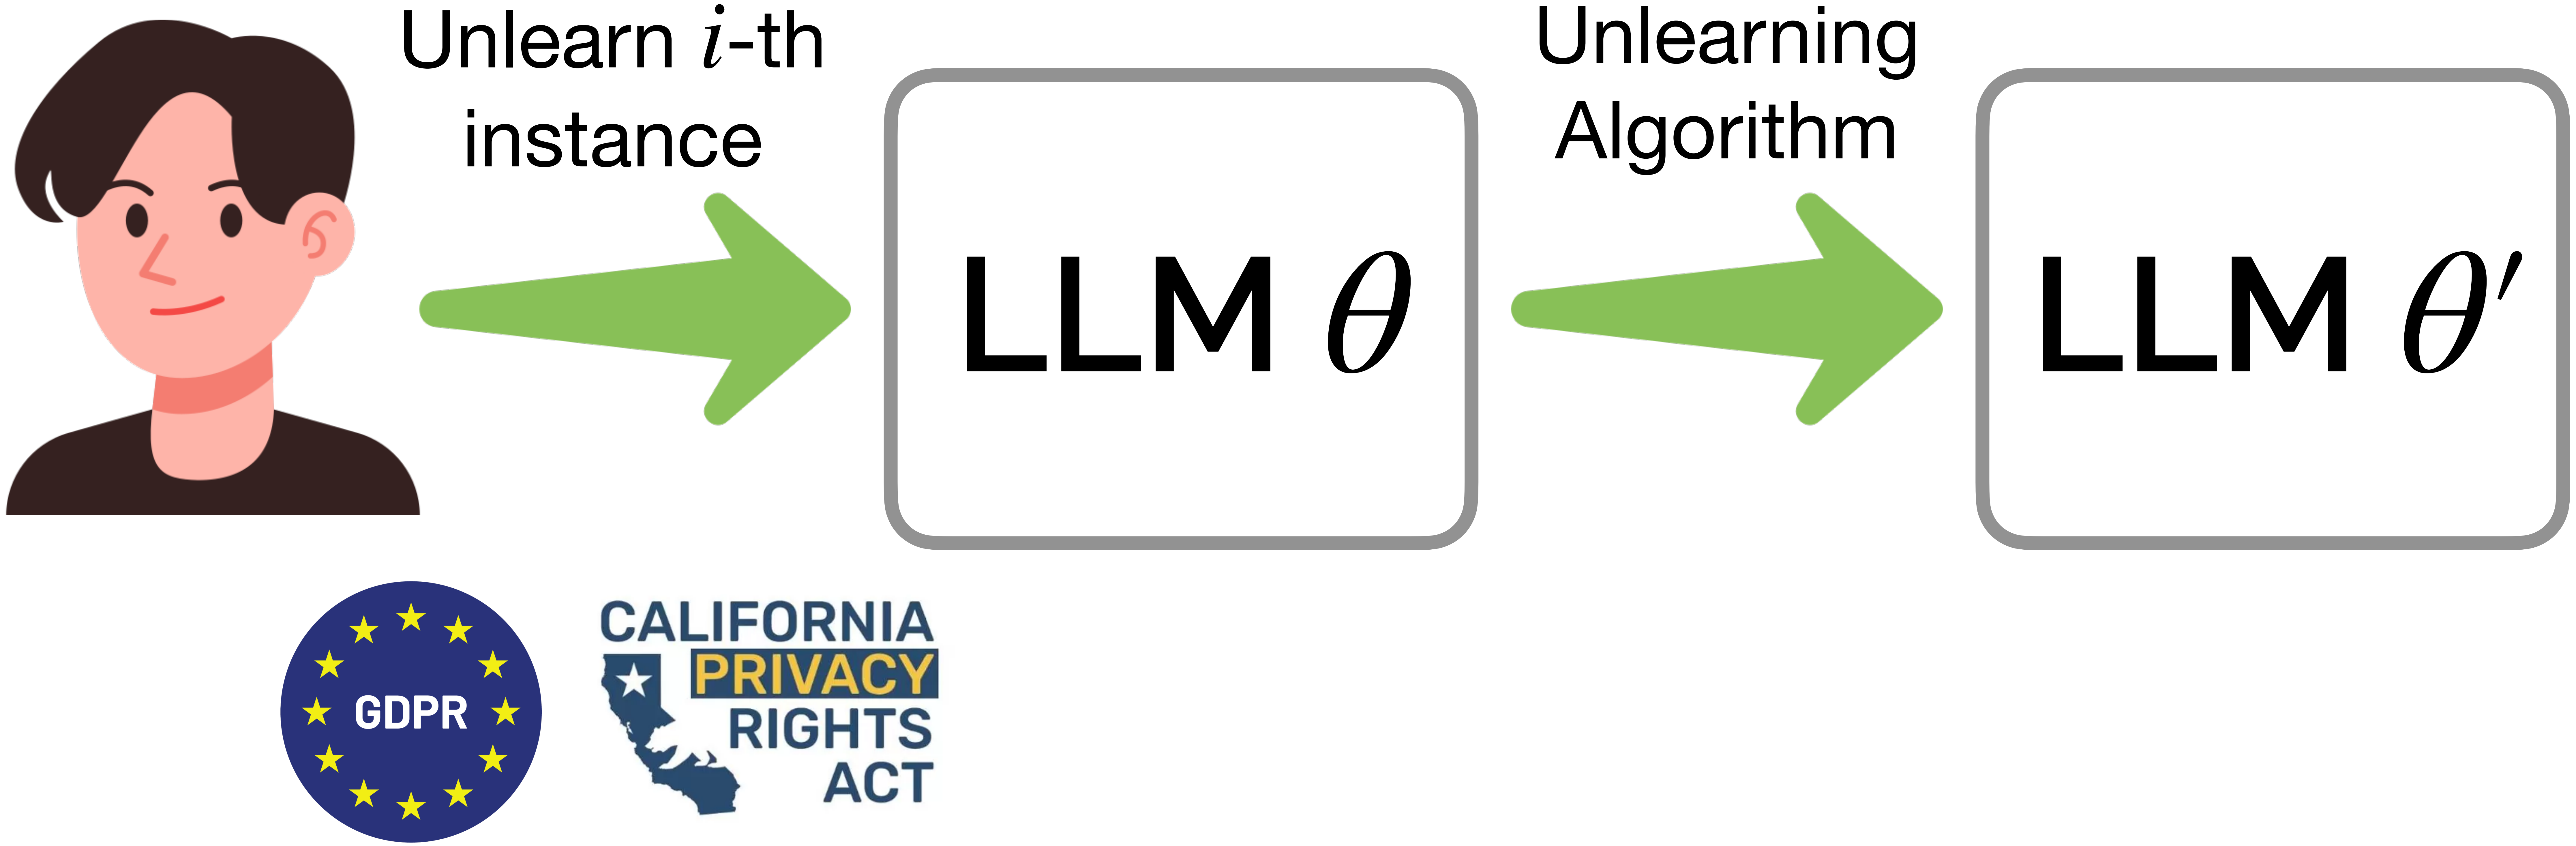

Unlearn i-th instance

LLM θ

Unlearning Algorithm

LLM θ'

Exact Unlearning

Exact unlearning guarantees that the ML model has perfectly erased information.

Exact Unlearning: Modular System

Machine
Learning Model

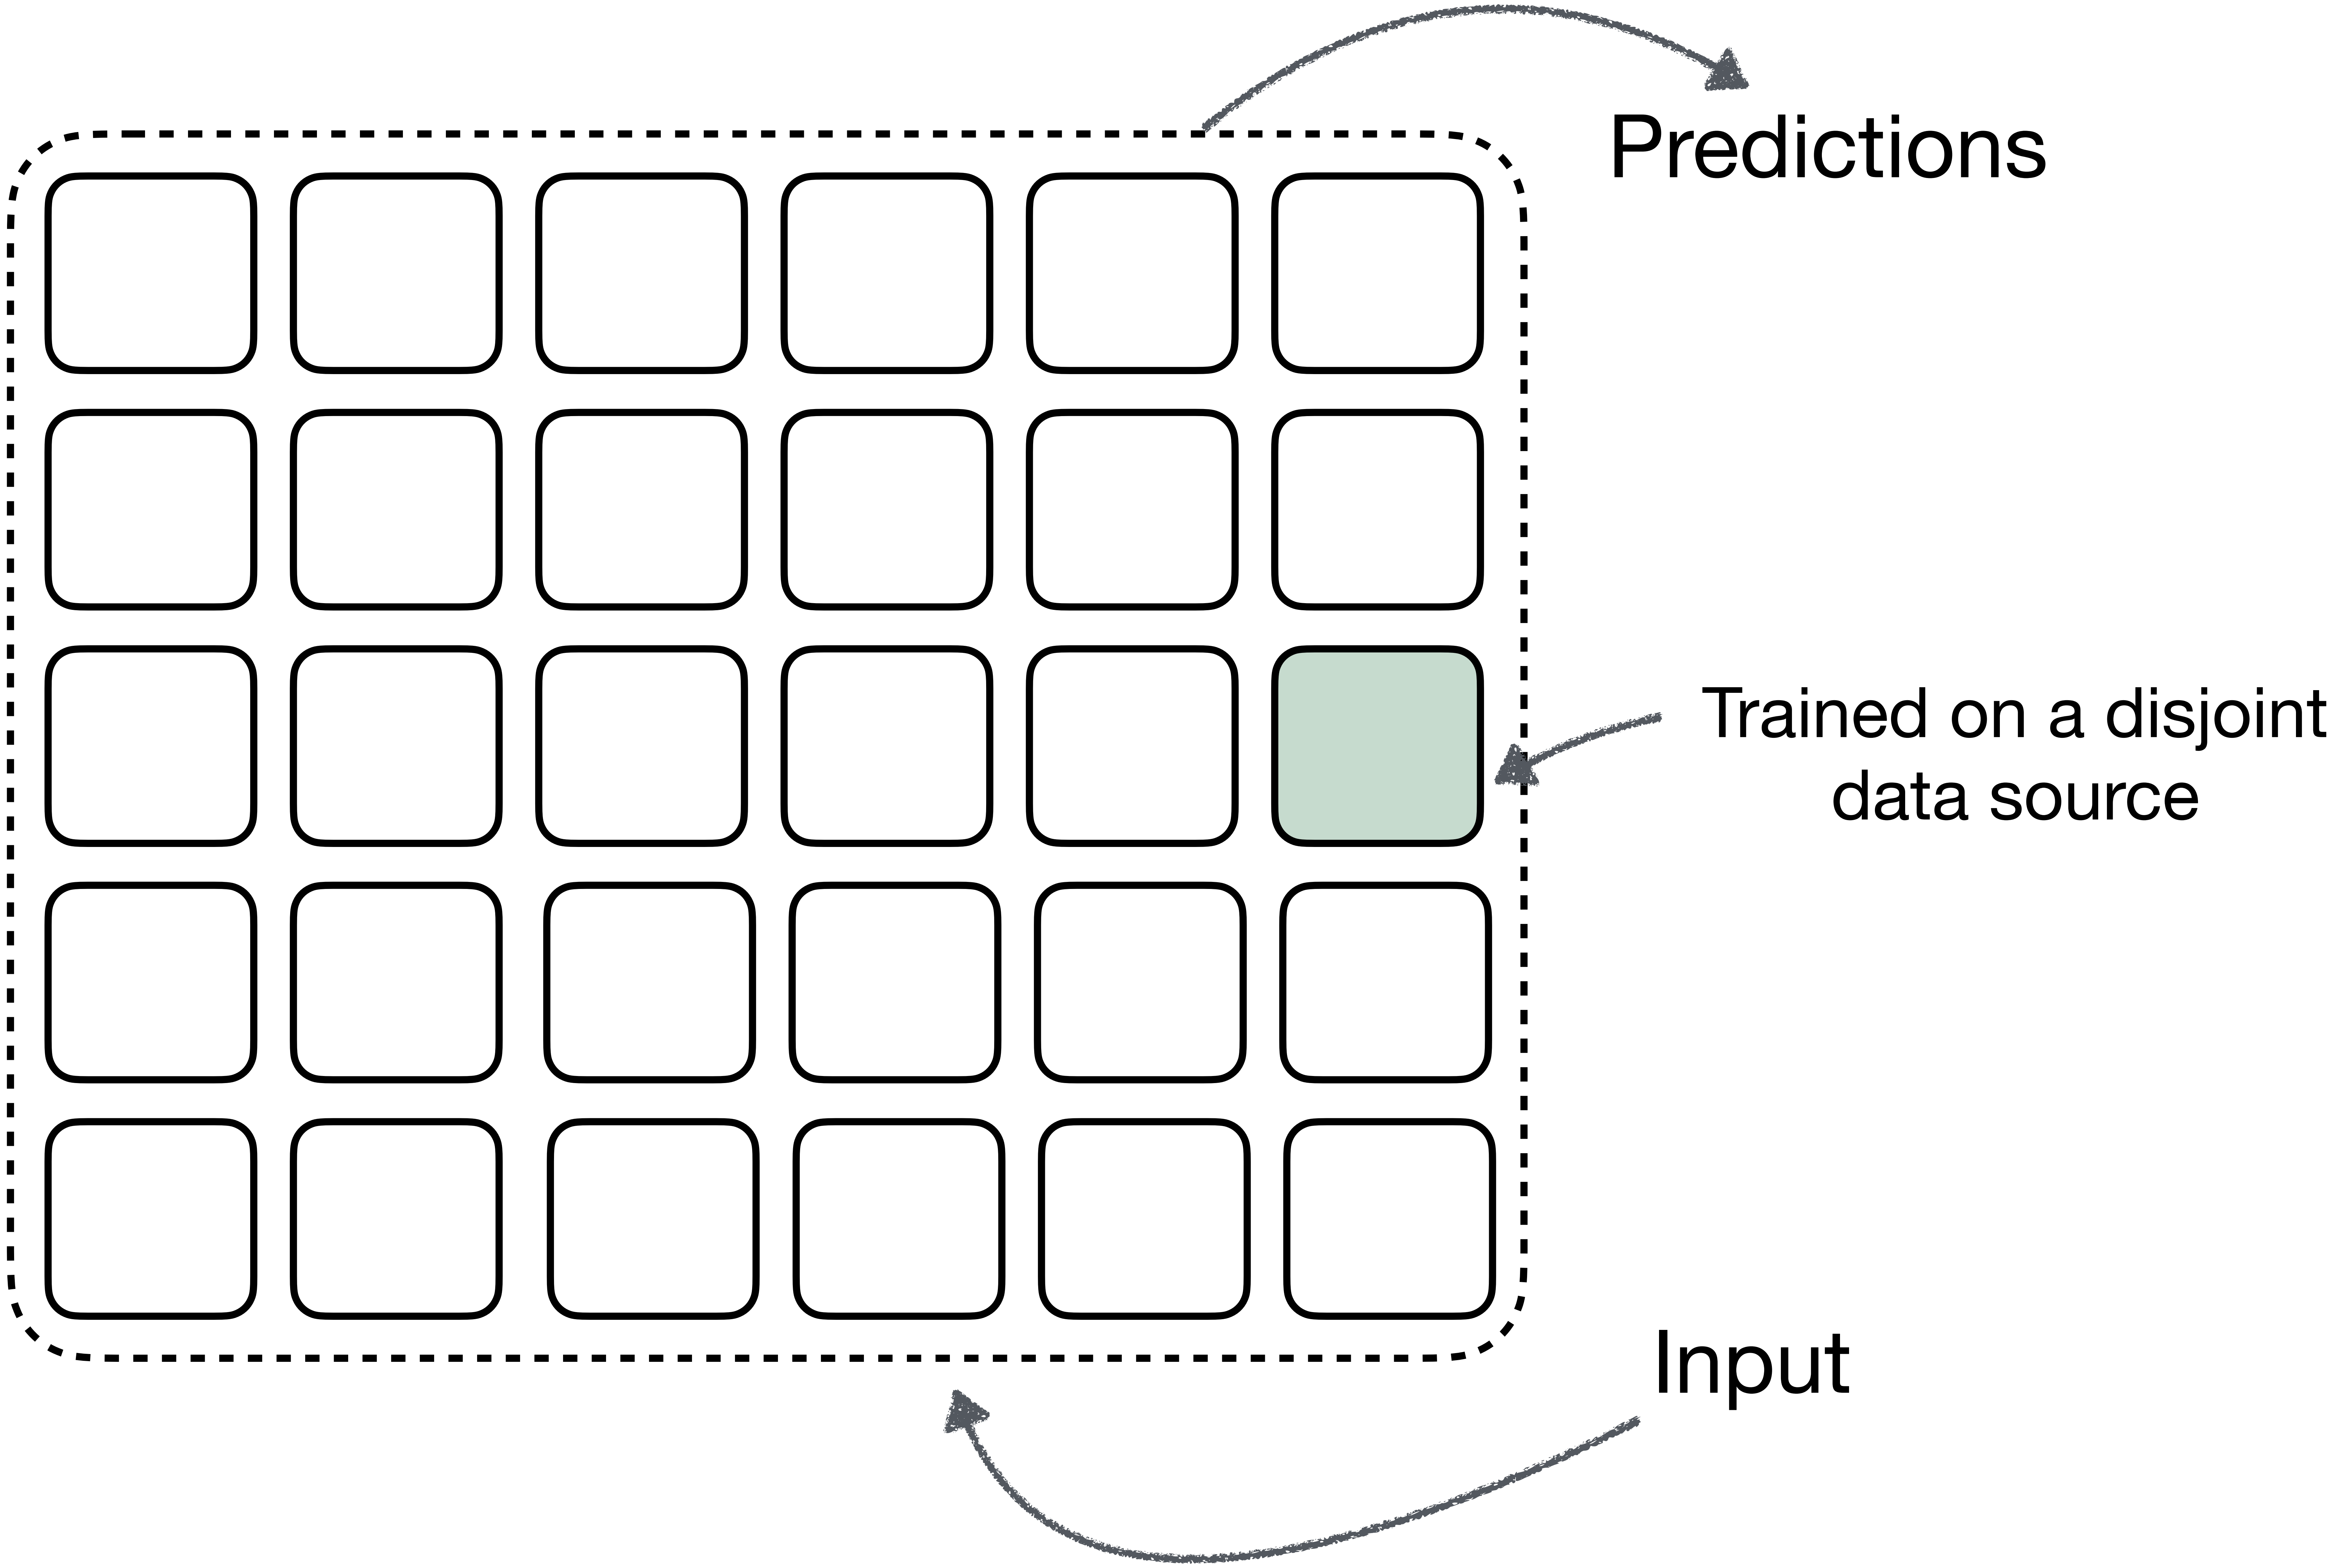

Predictions

Trained on a disjoint data source

Input

Exact Unlearning: Modular System

Machine
Learning Model

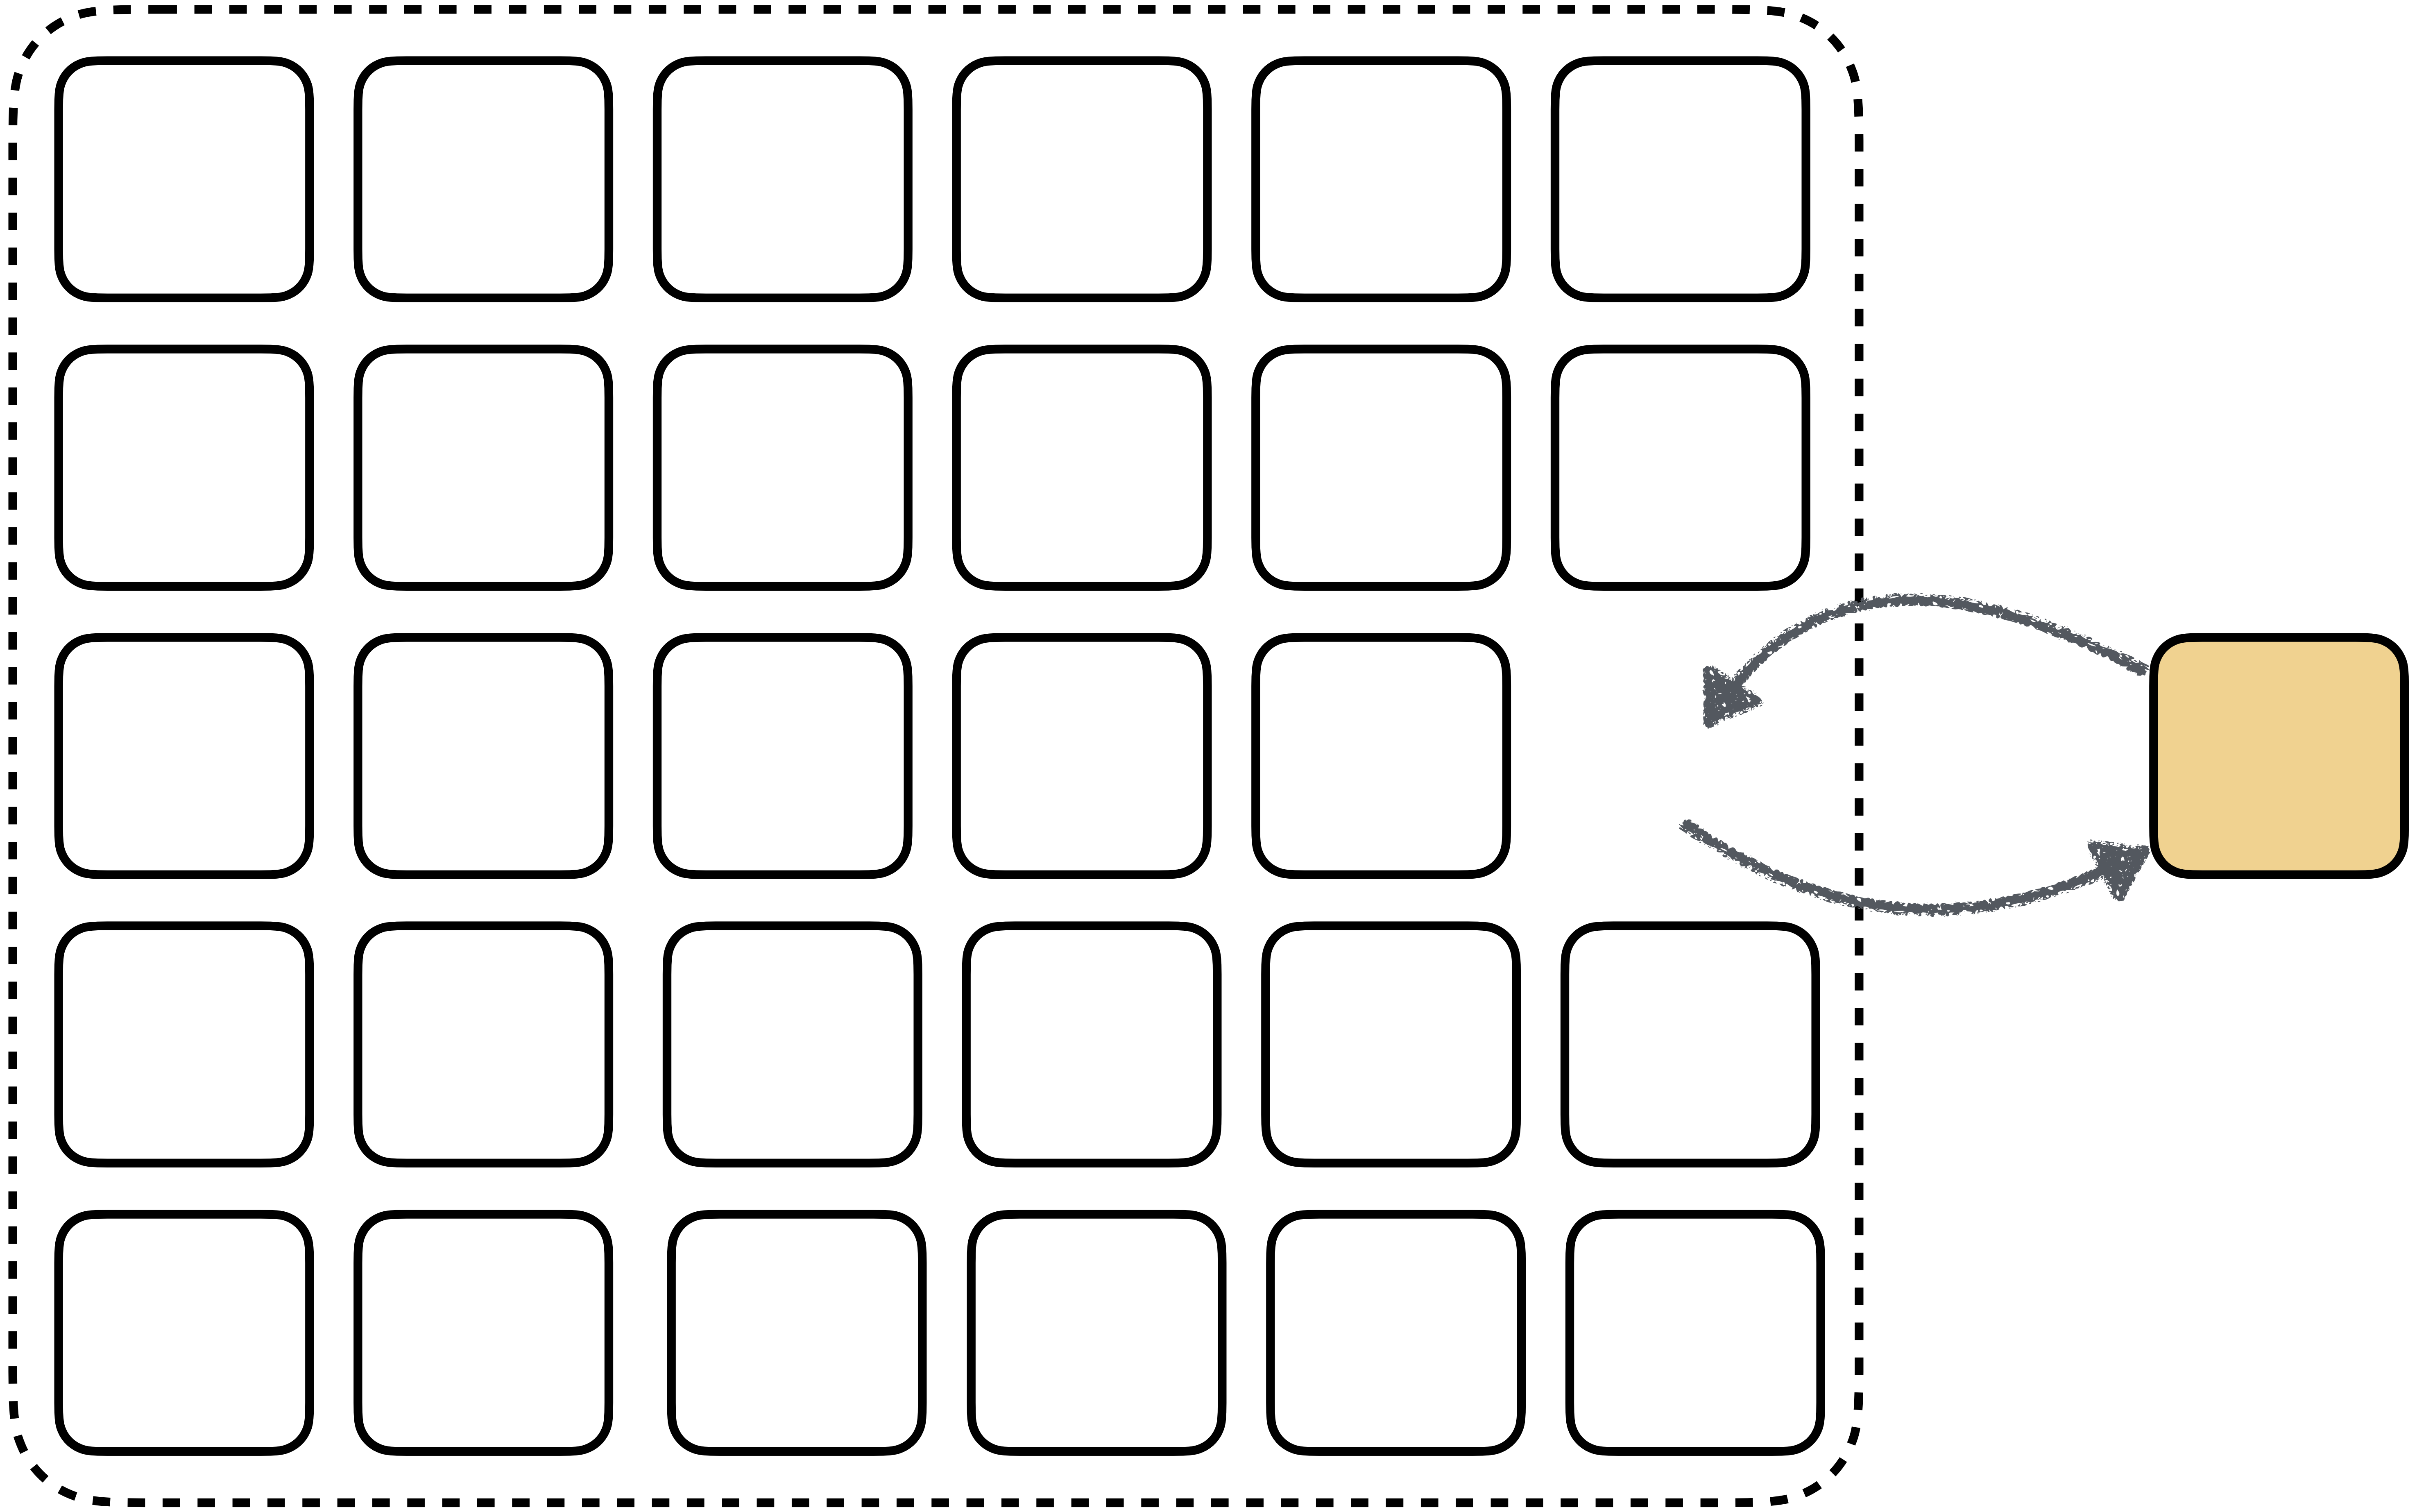

Retrain on
$1/N$-th data

S^{3}T: Sequential Slice-aware Training

S^3T: Sequential Slice-aware Training

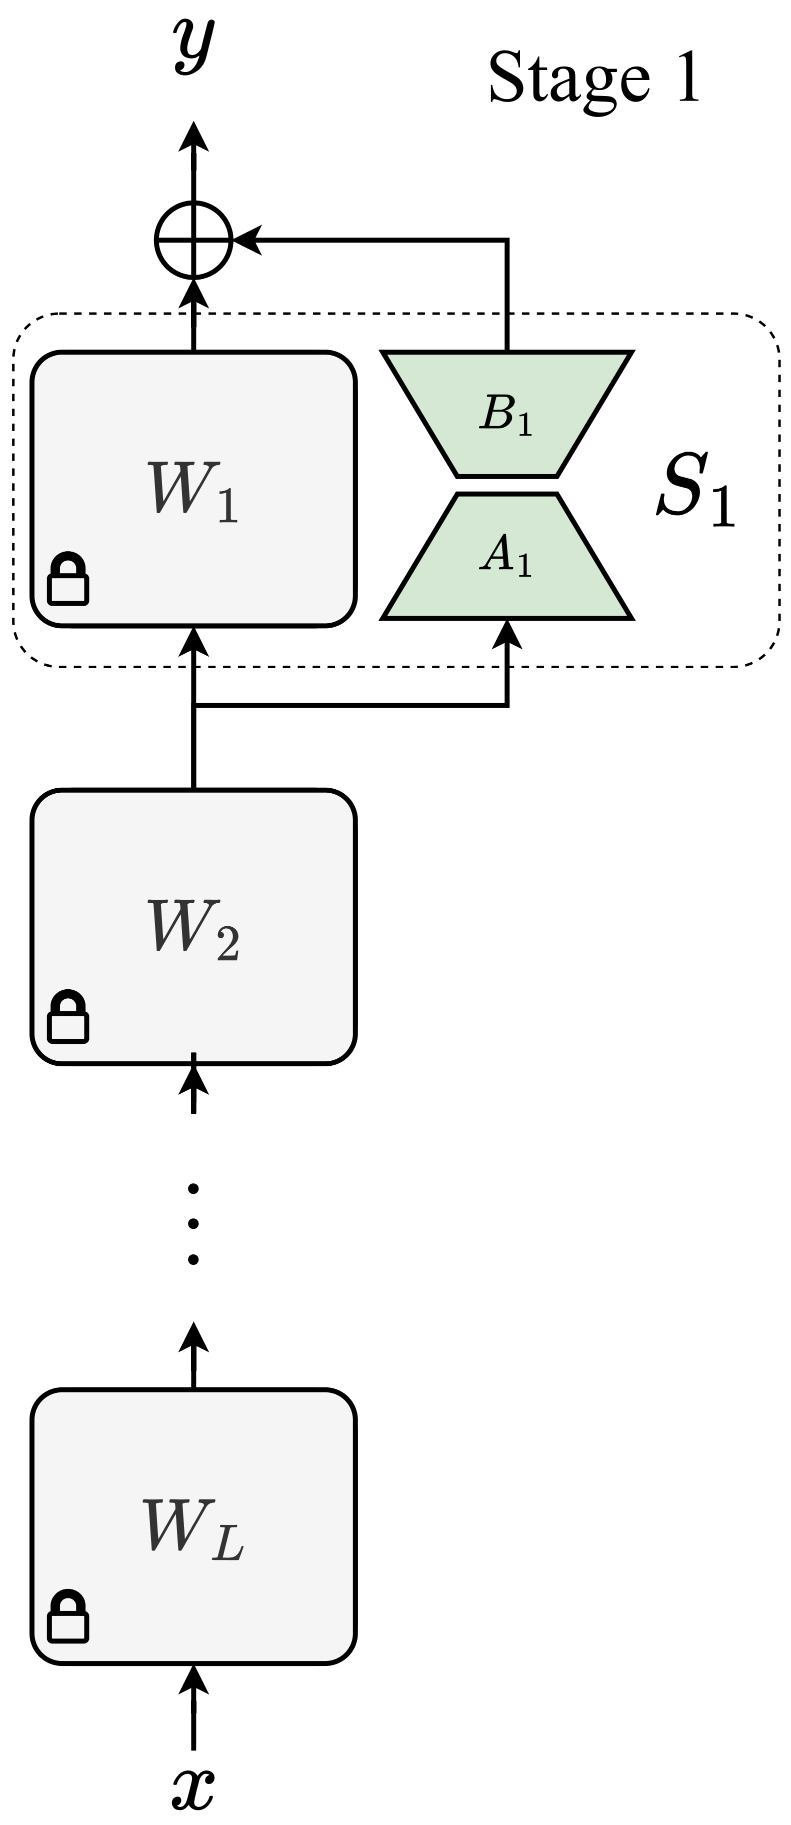

S^{3}T: Sequential Slice-aware Training

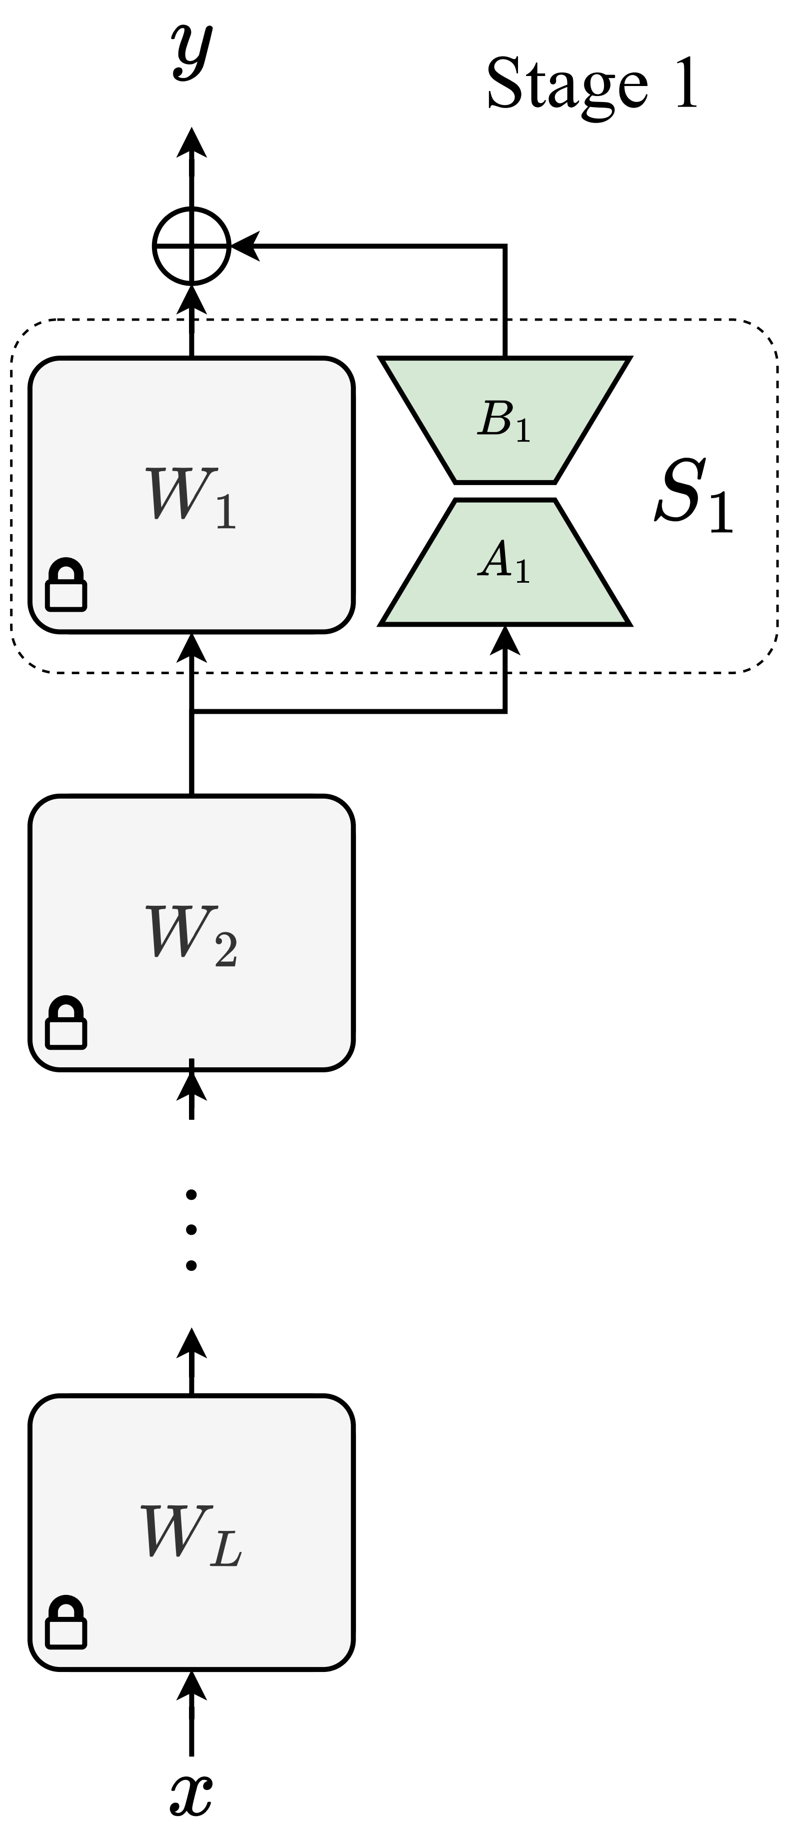

Train the top-layer with slice S_1 while freezing the rest of the layers.

S^3T: Sequential Slice-aware Training

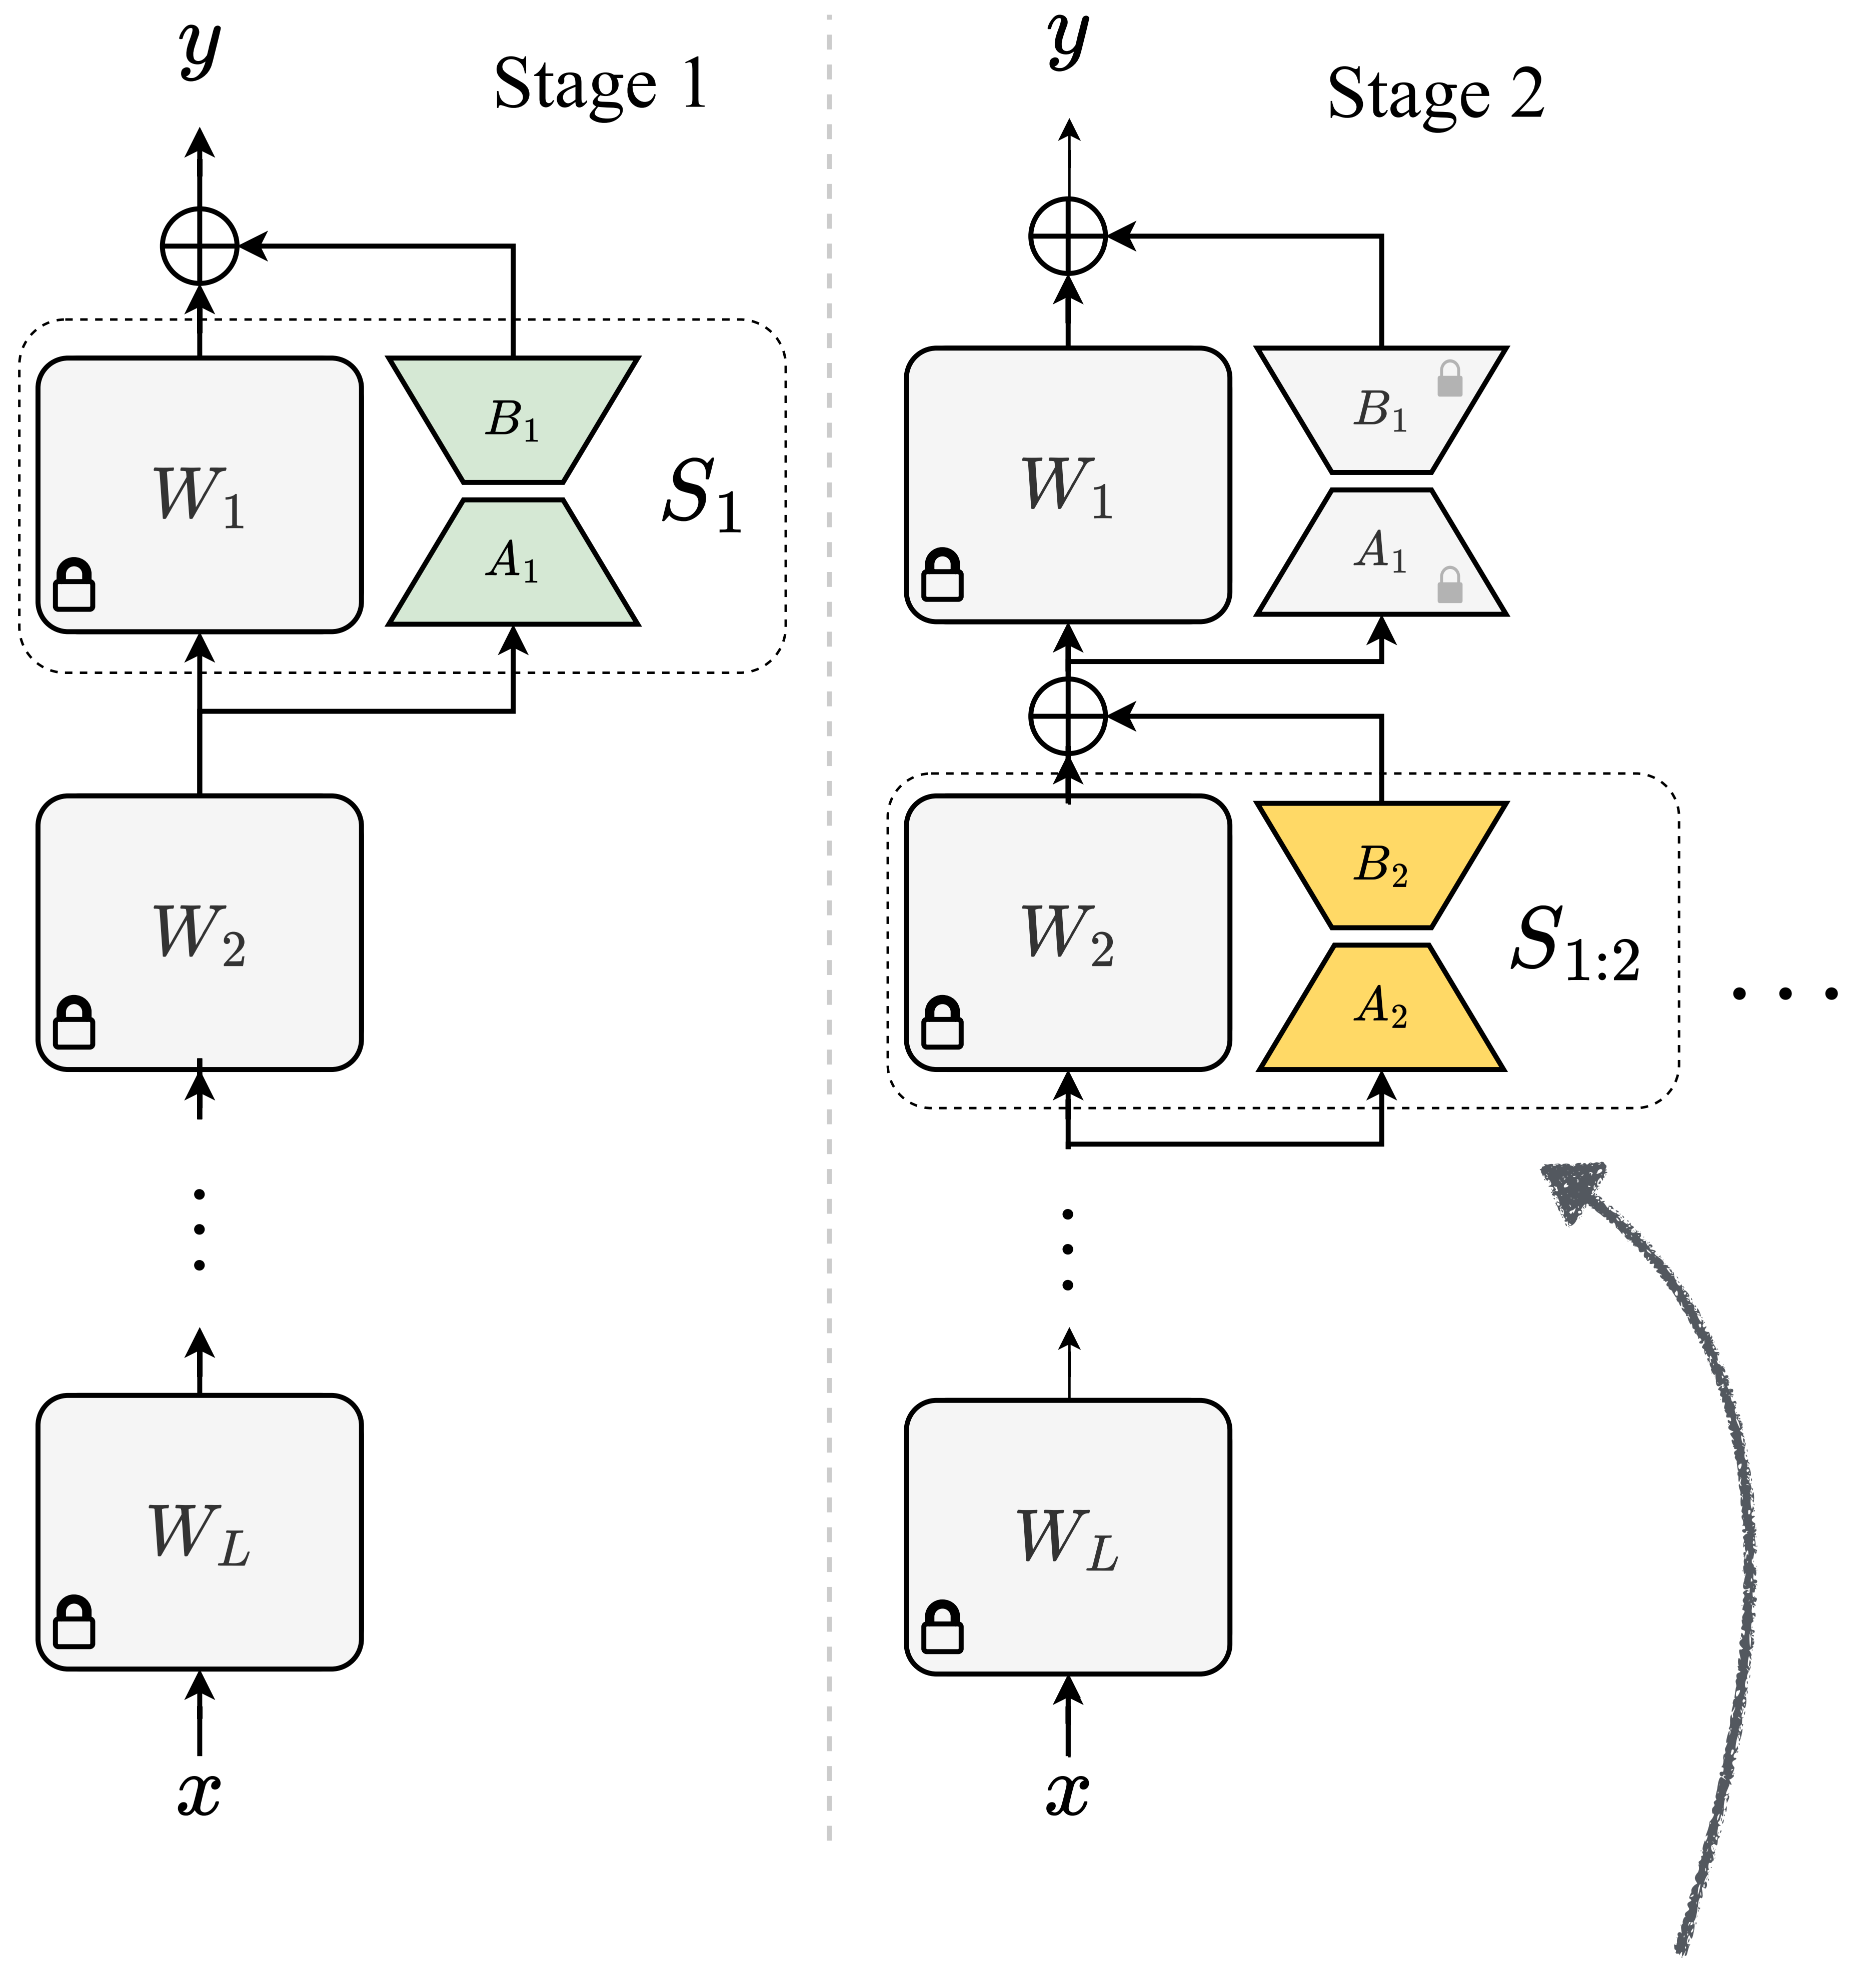

The second layer is trained using slices ($S_1 + S_2$).

S^{3}T: Sequential Slice-aware Training

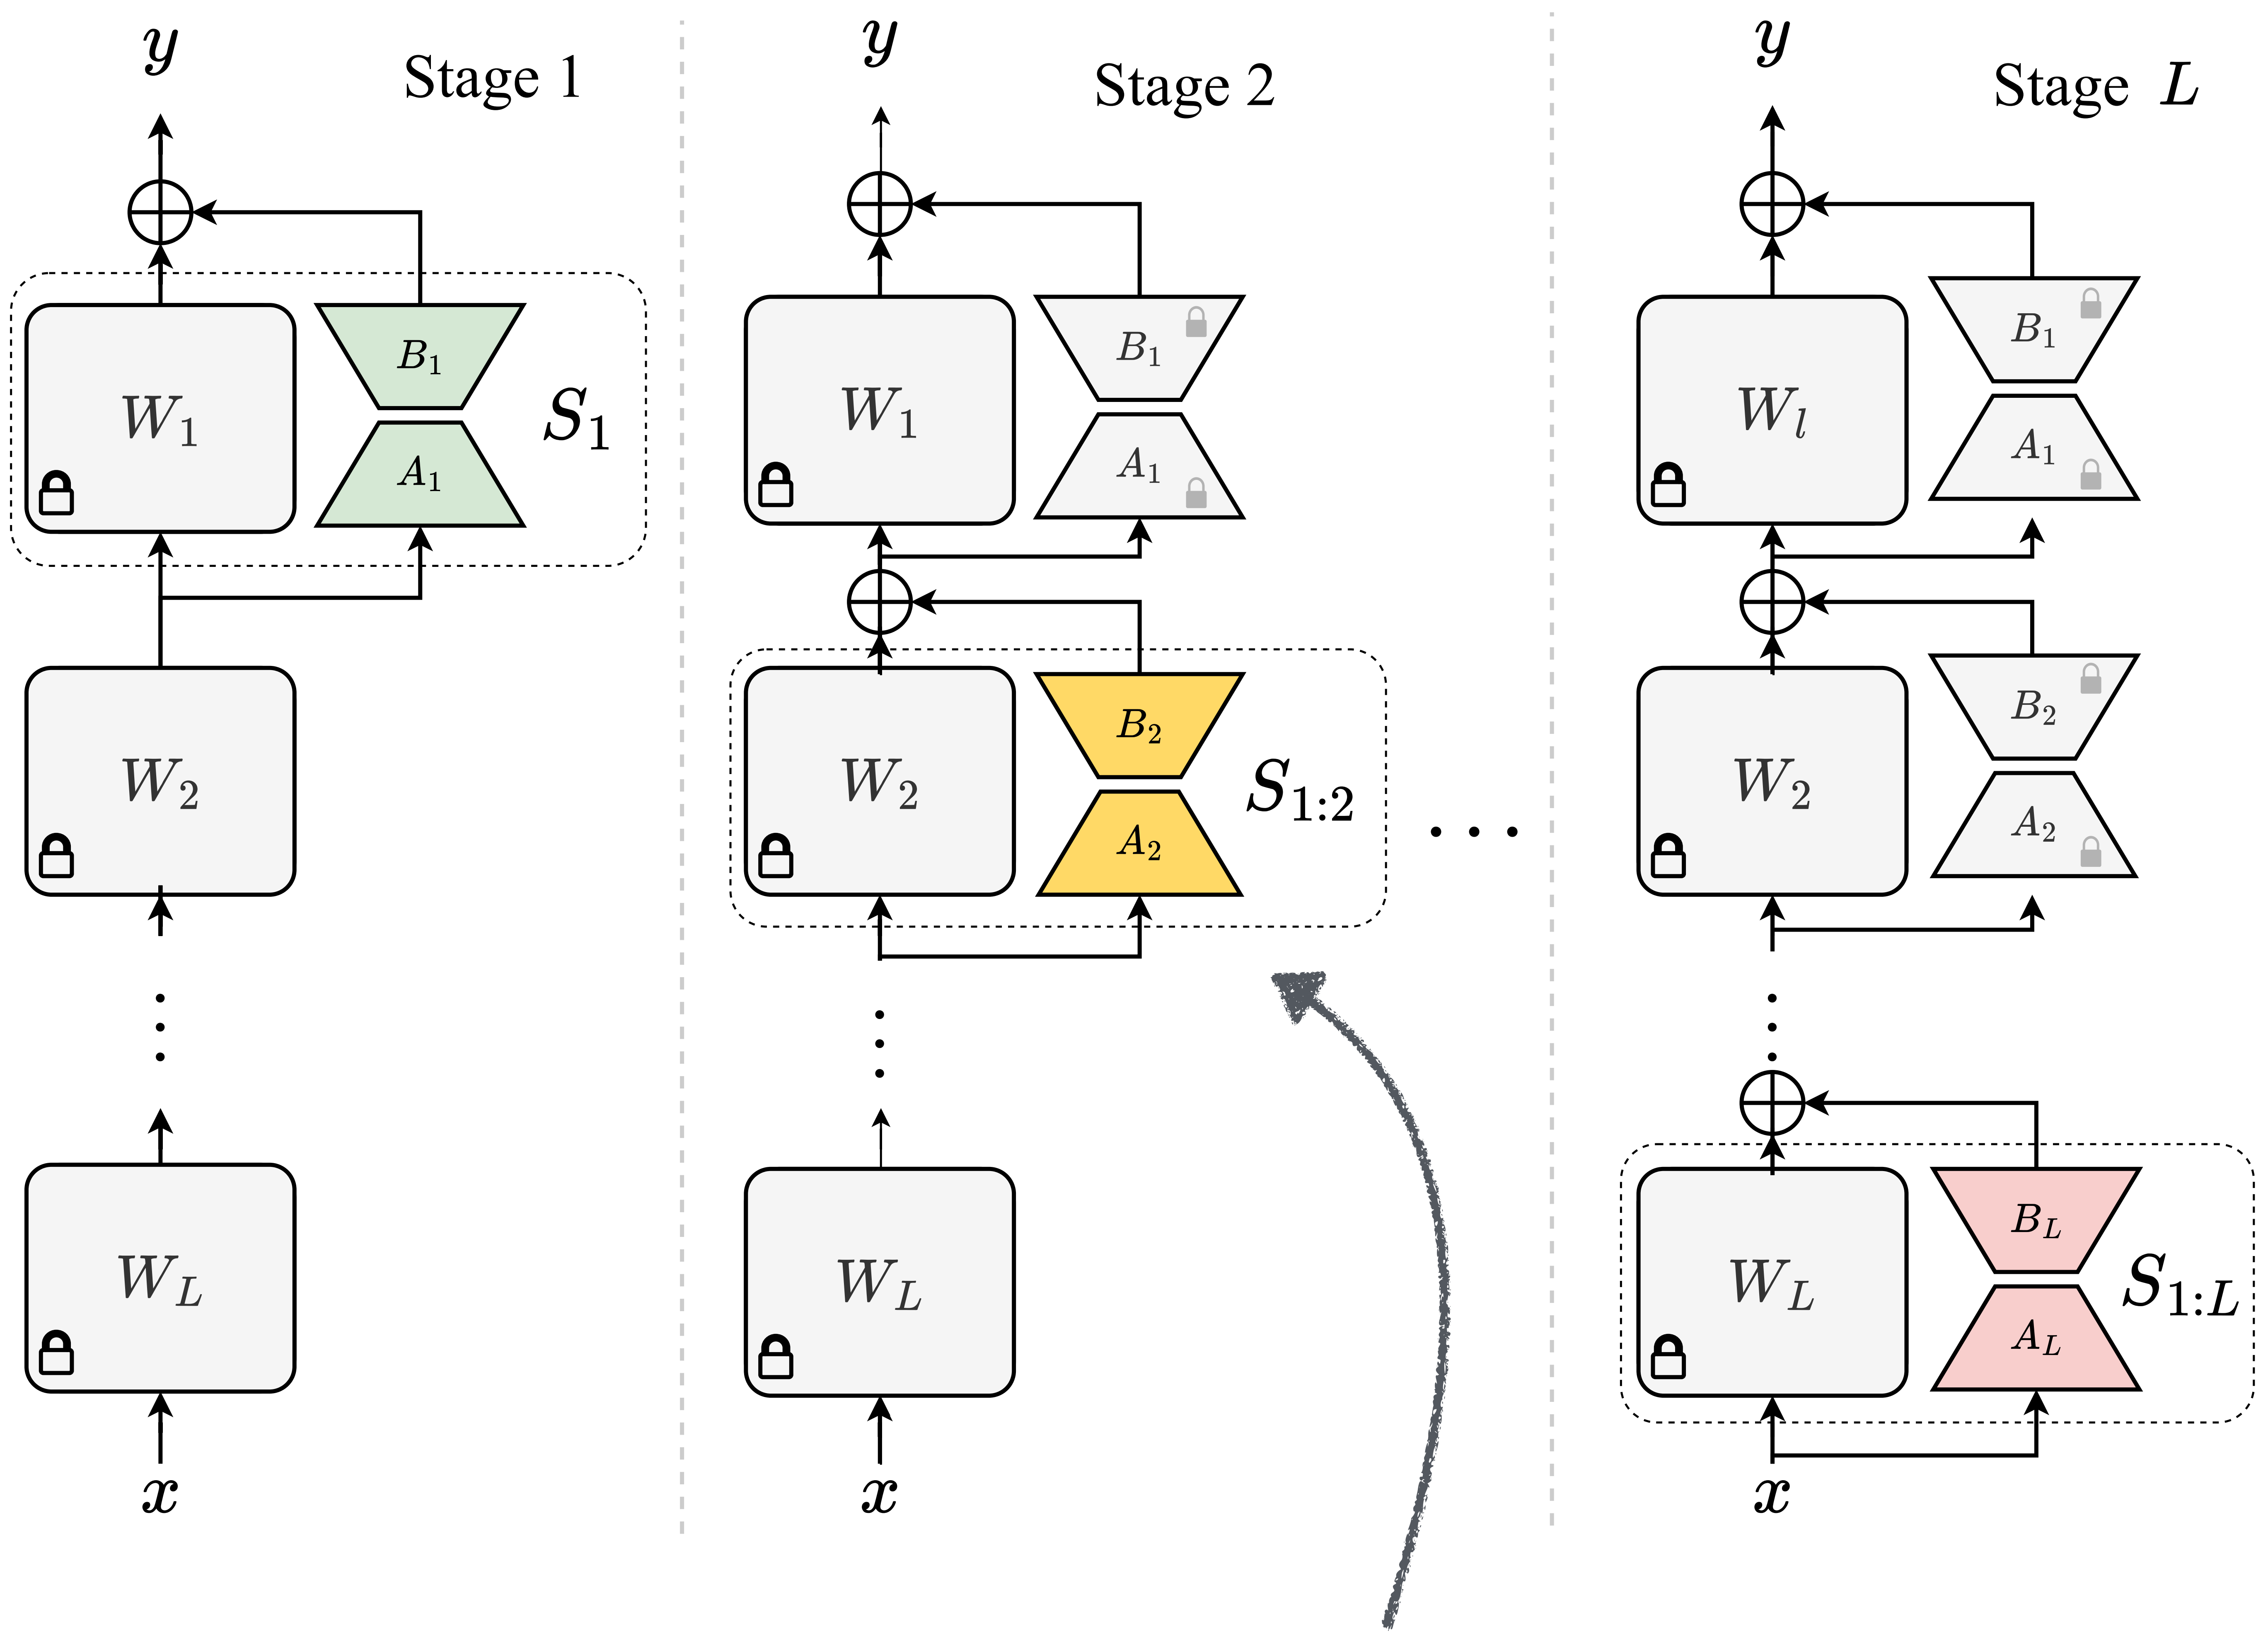

The second layer is trained using slices ($S_1 + S_2$). This continues.

S^3T: Sequential Slice-aware Training

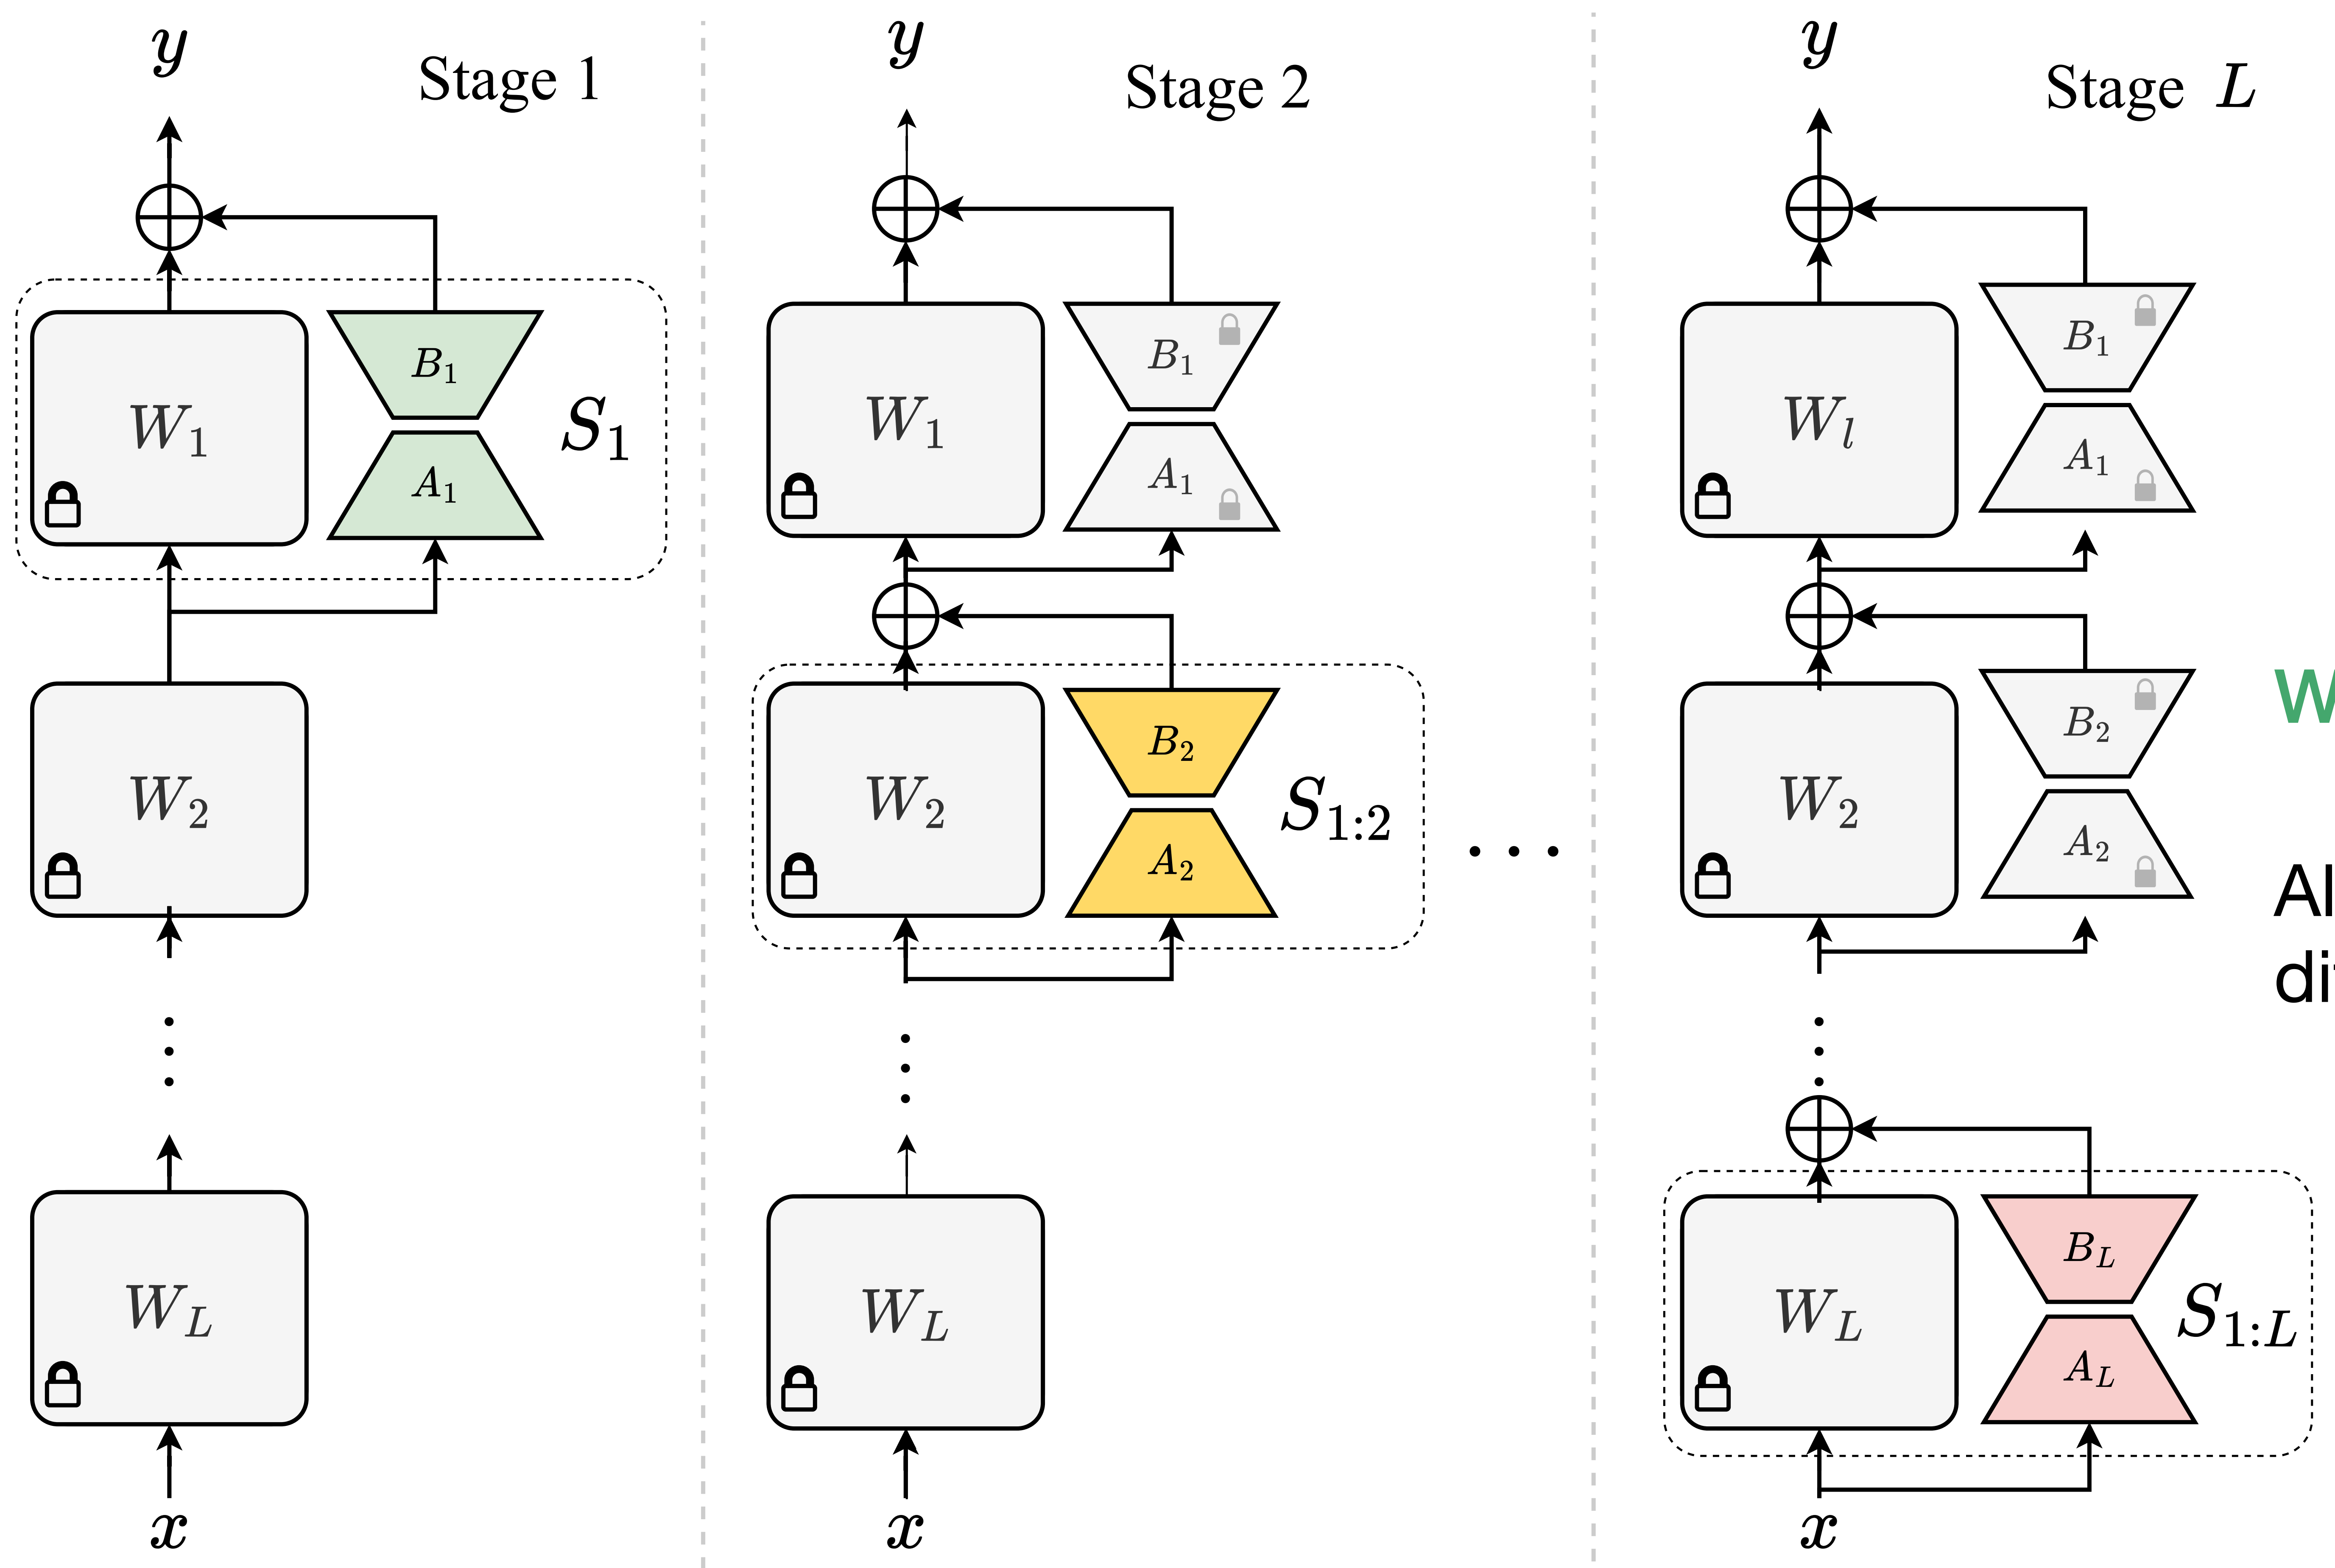

Why do we even need this?

Allows parameter isolation for different slices.

S^{3}T: Sequential Slice-aware Training

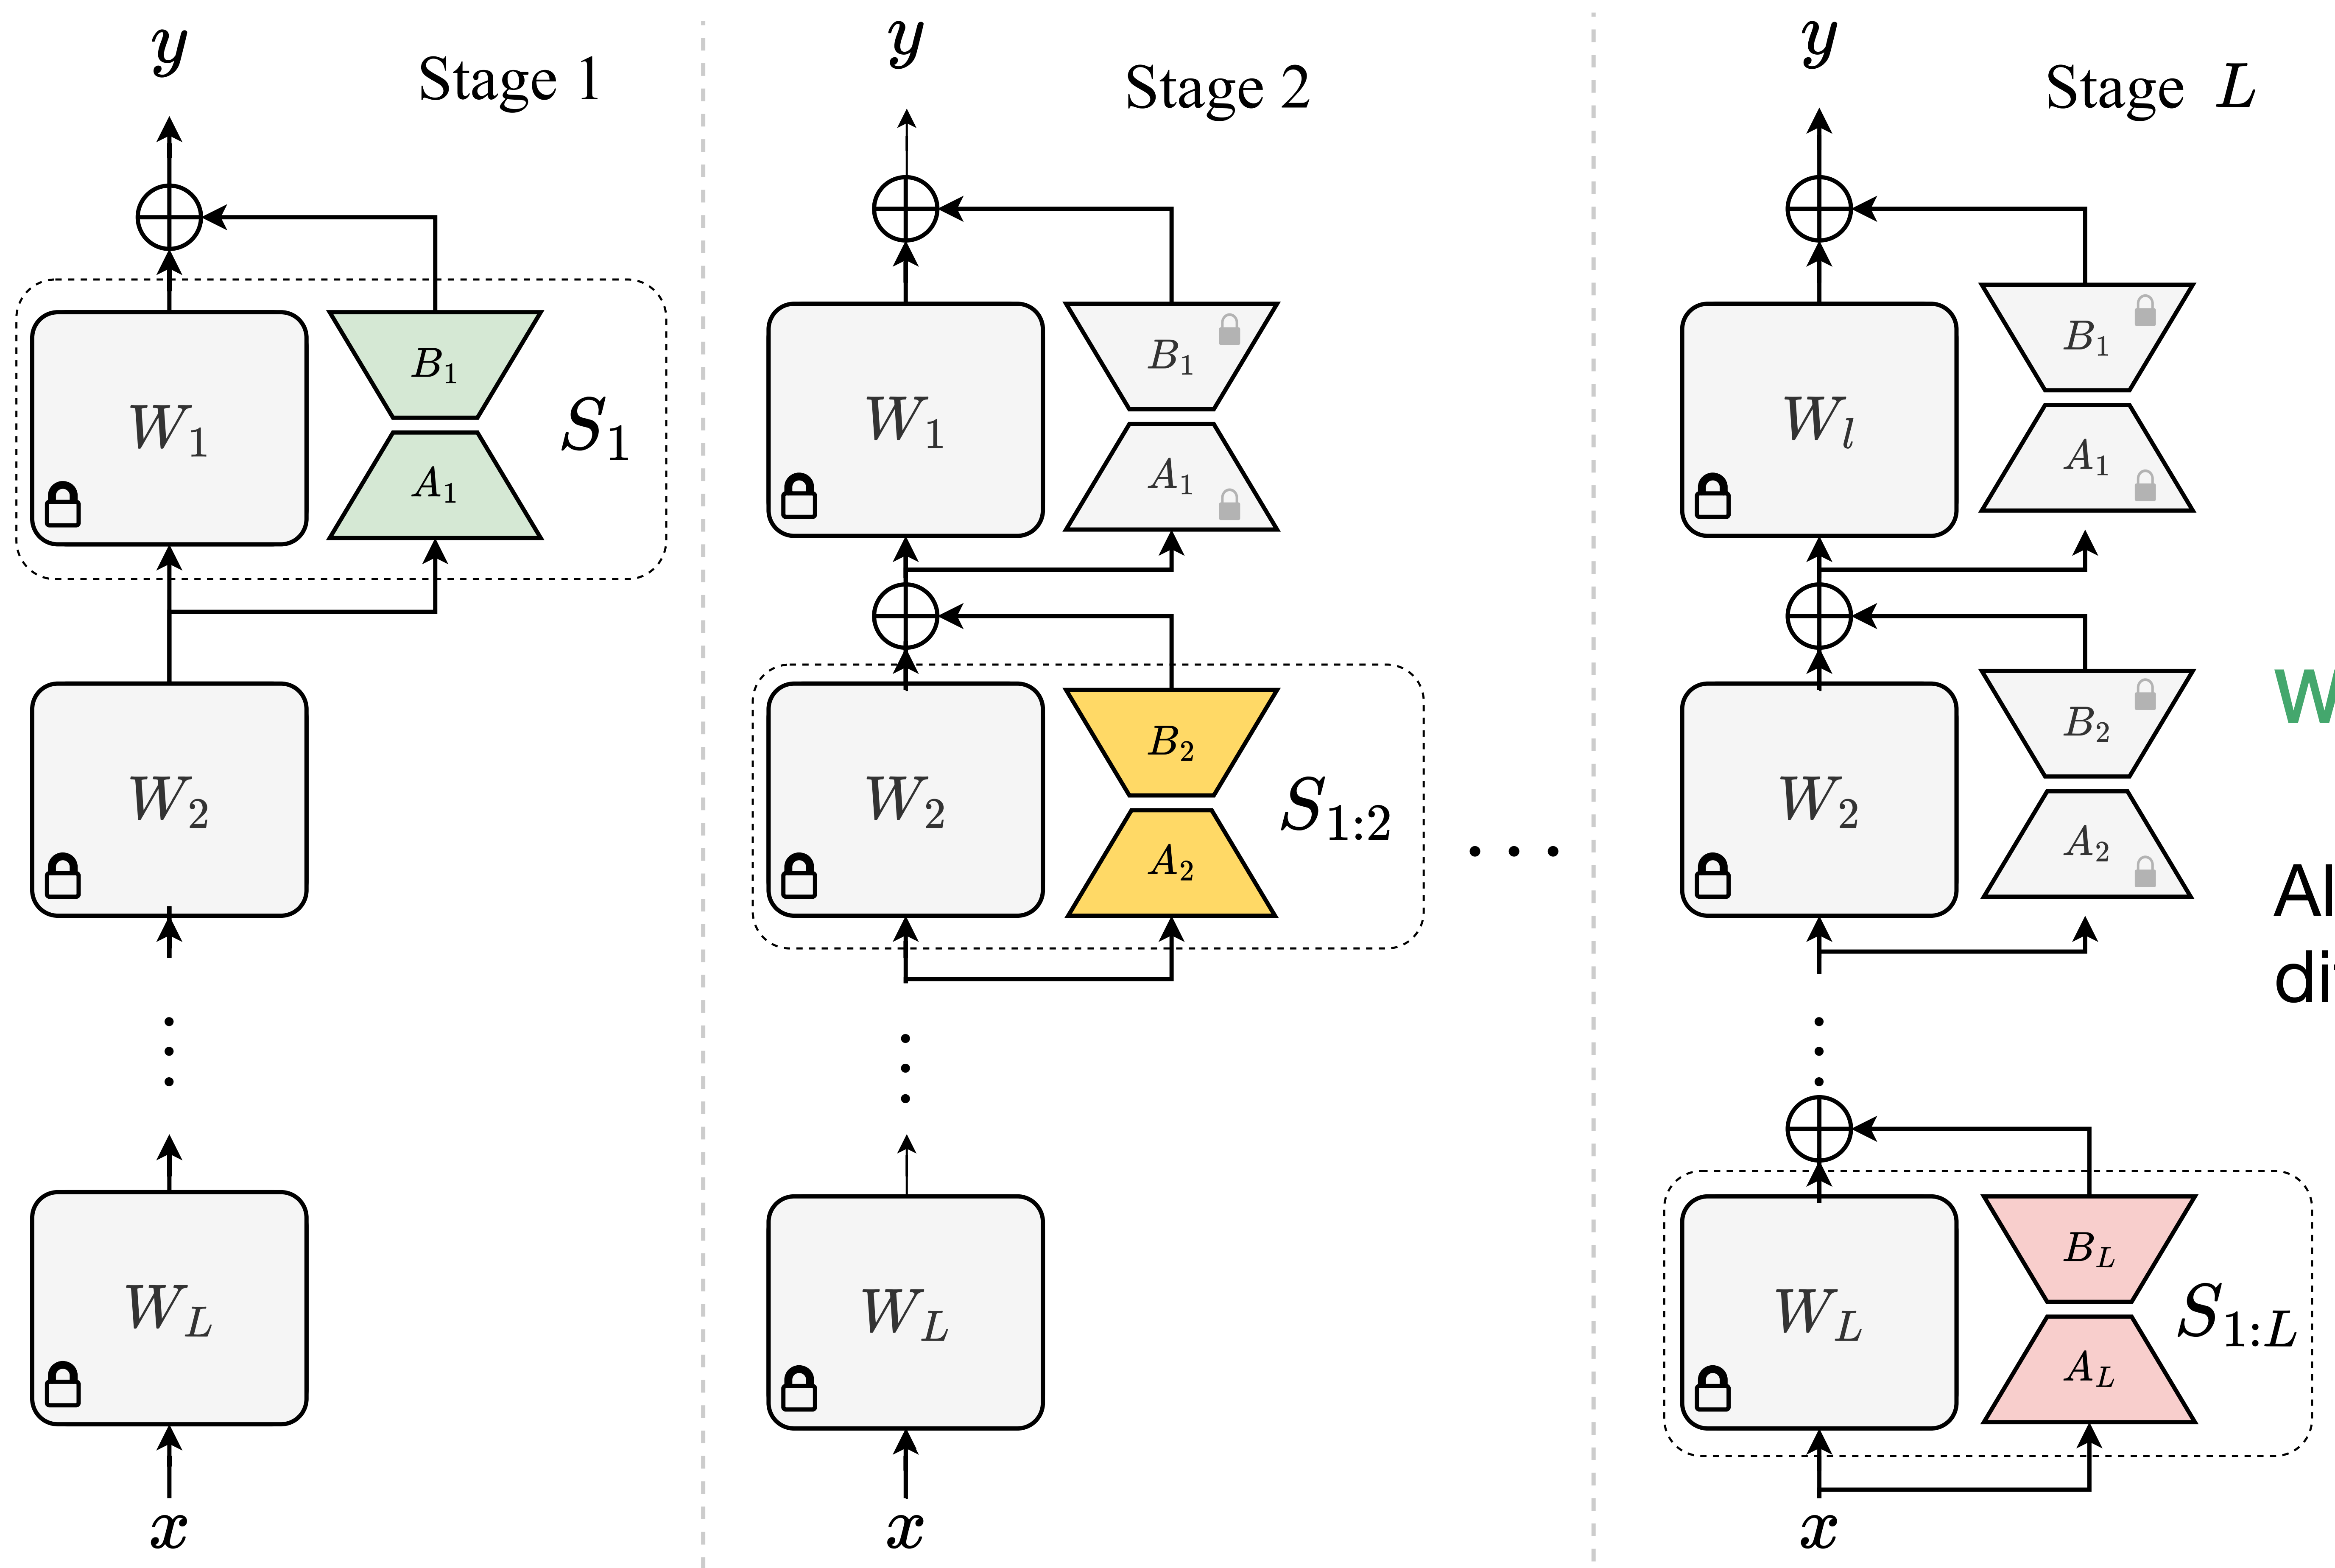

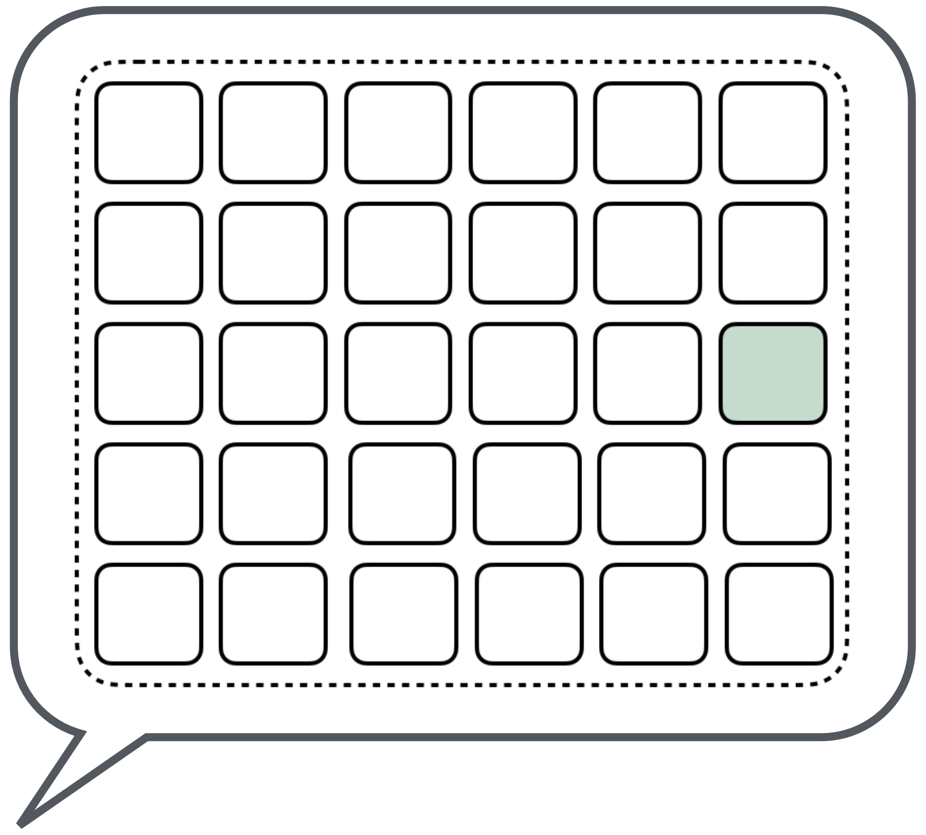

Why do we even need this?

Allows parameter isolation for different slices.

S^{3}T: Sequential Slice-aware Training

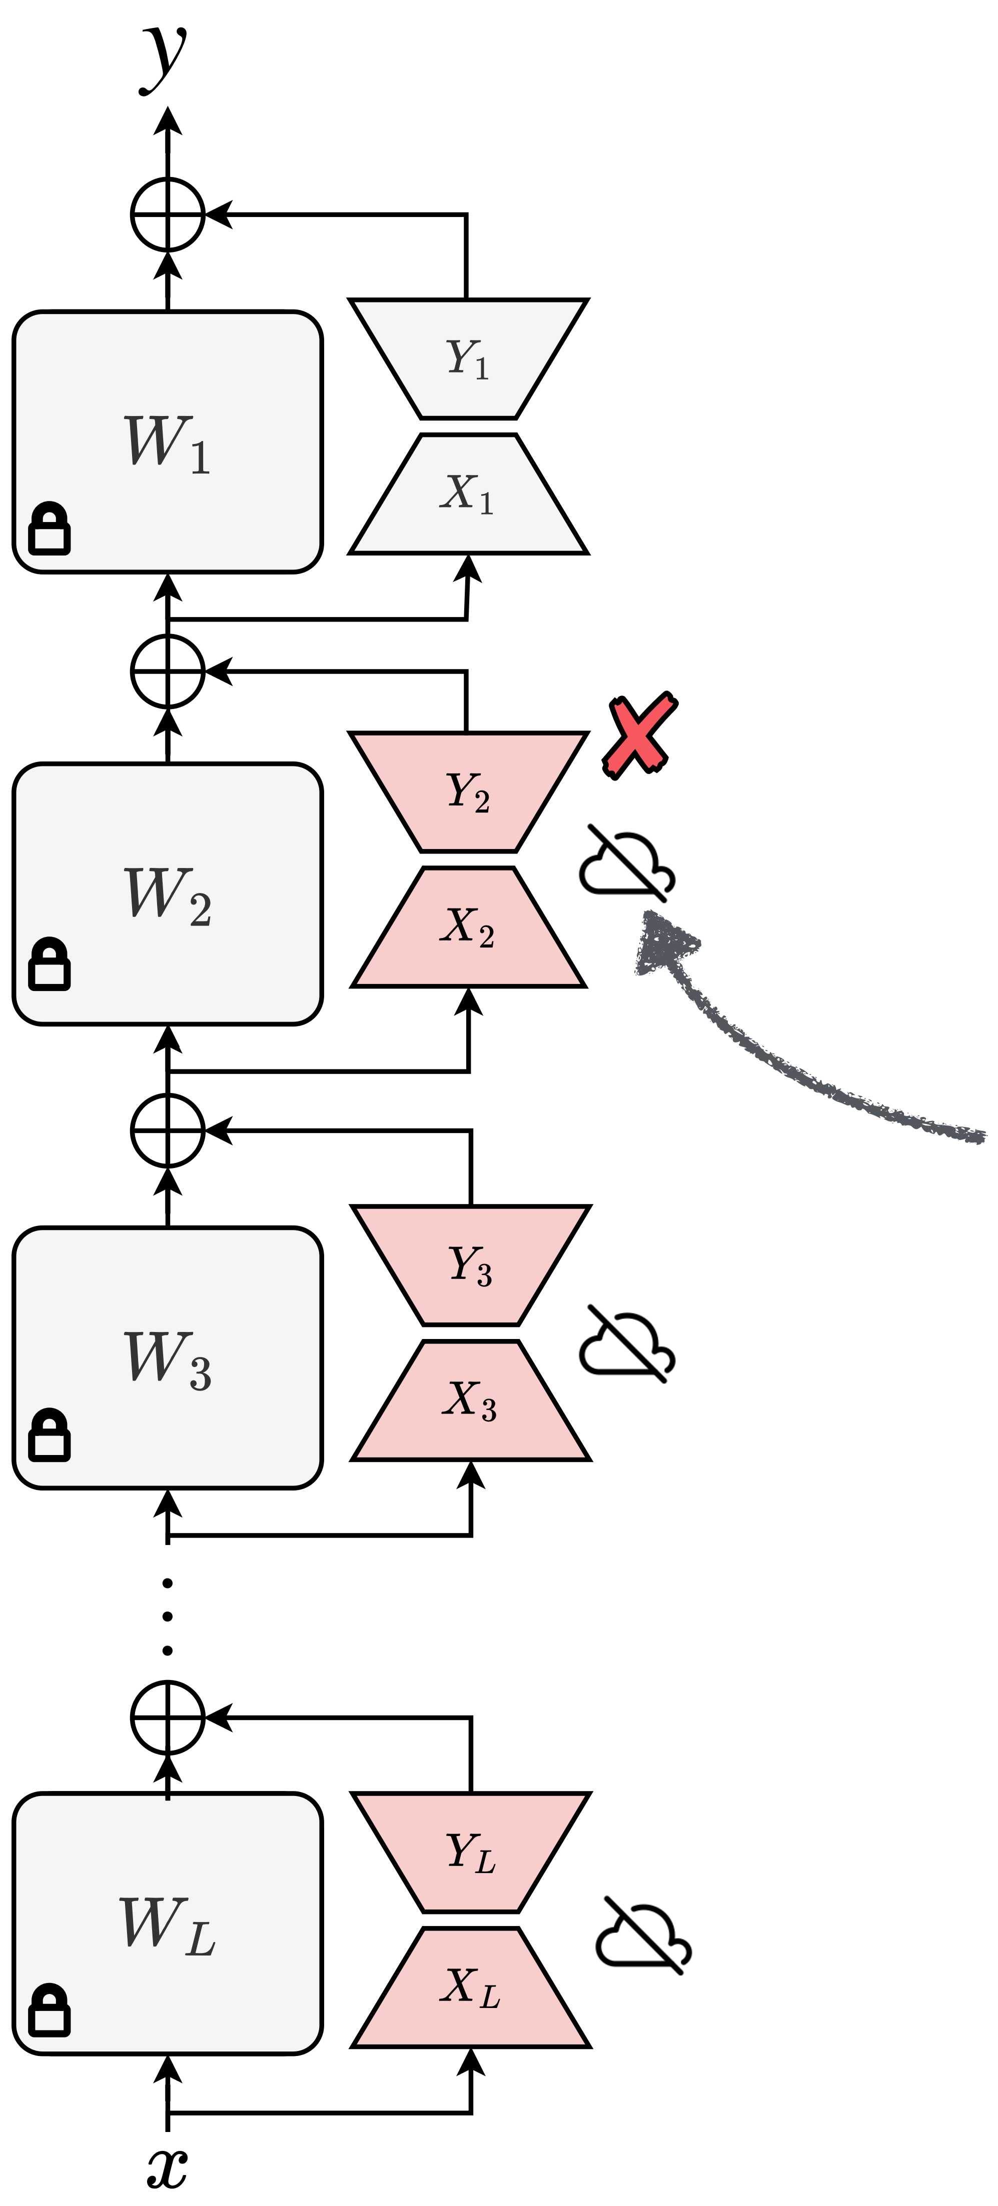

If a deletion request affects S_2, it can be unlearned by switching off all PEFT layers below it

S^{3}T: Sequential Slice-aware Training

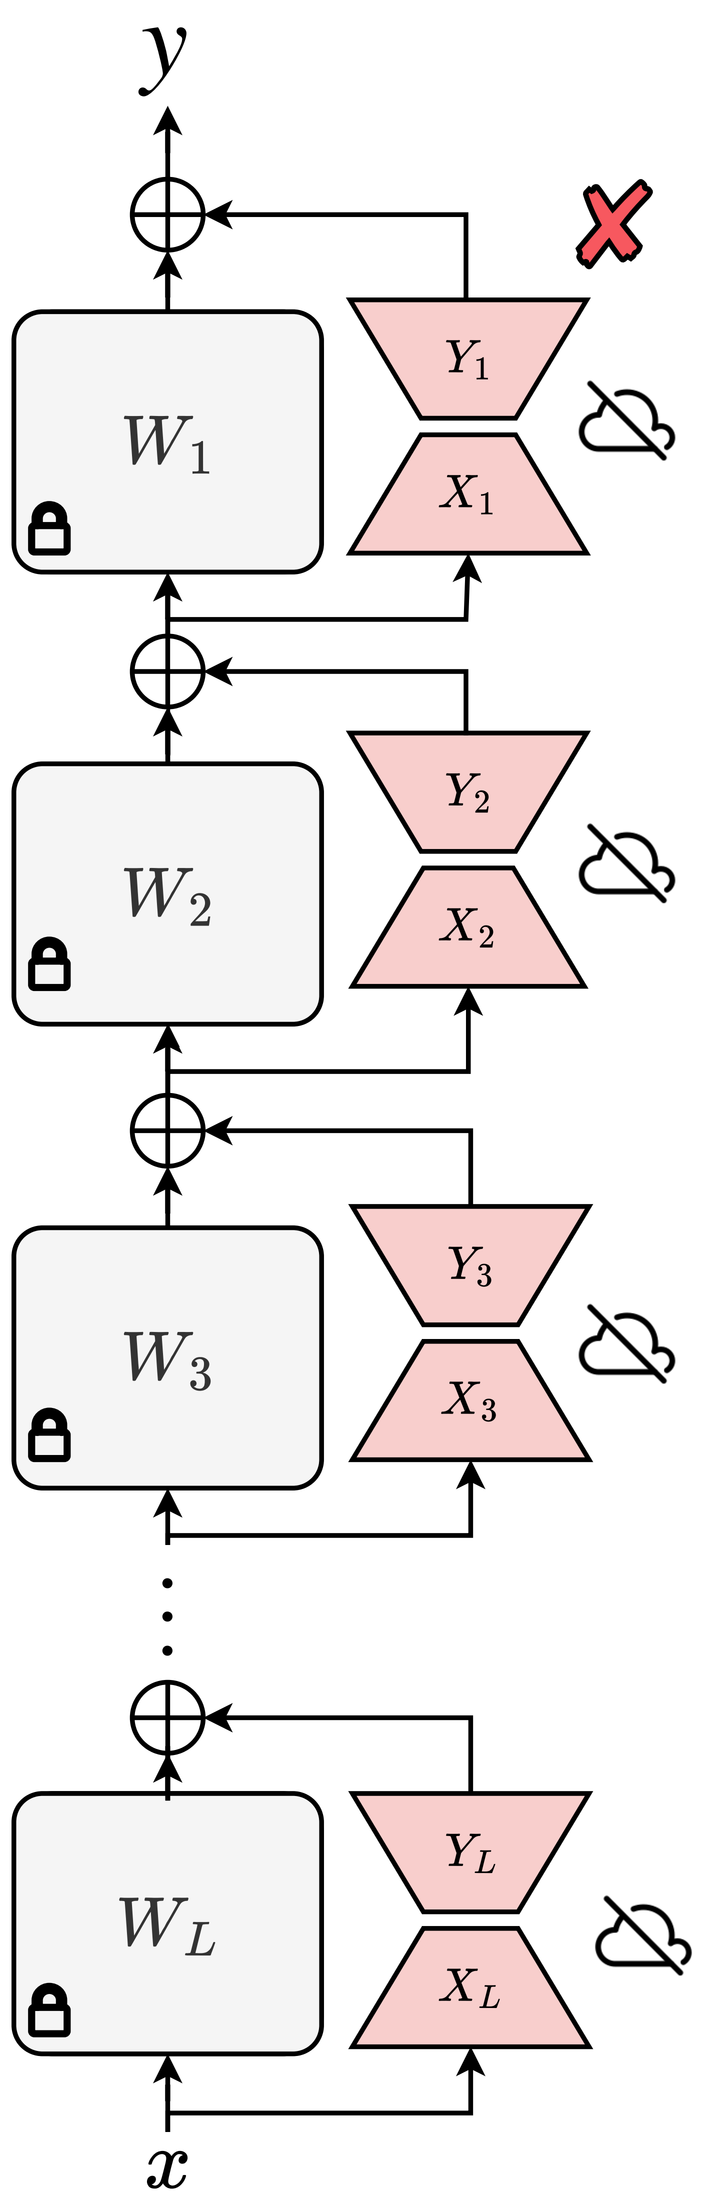

Switching off all PEFT layers -
Retrain from scratch.

S^{3}T: Sequence Selection

$L = 3, L! = 6$ sequences

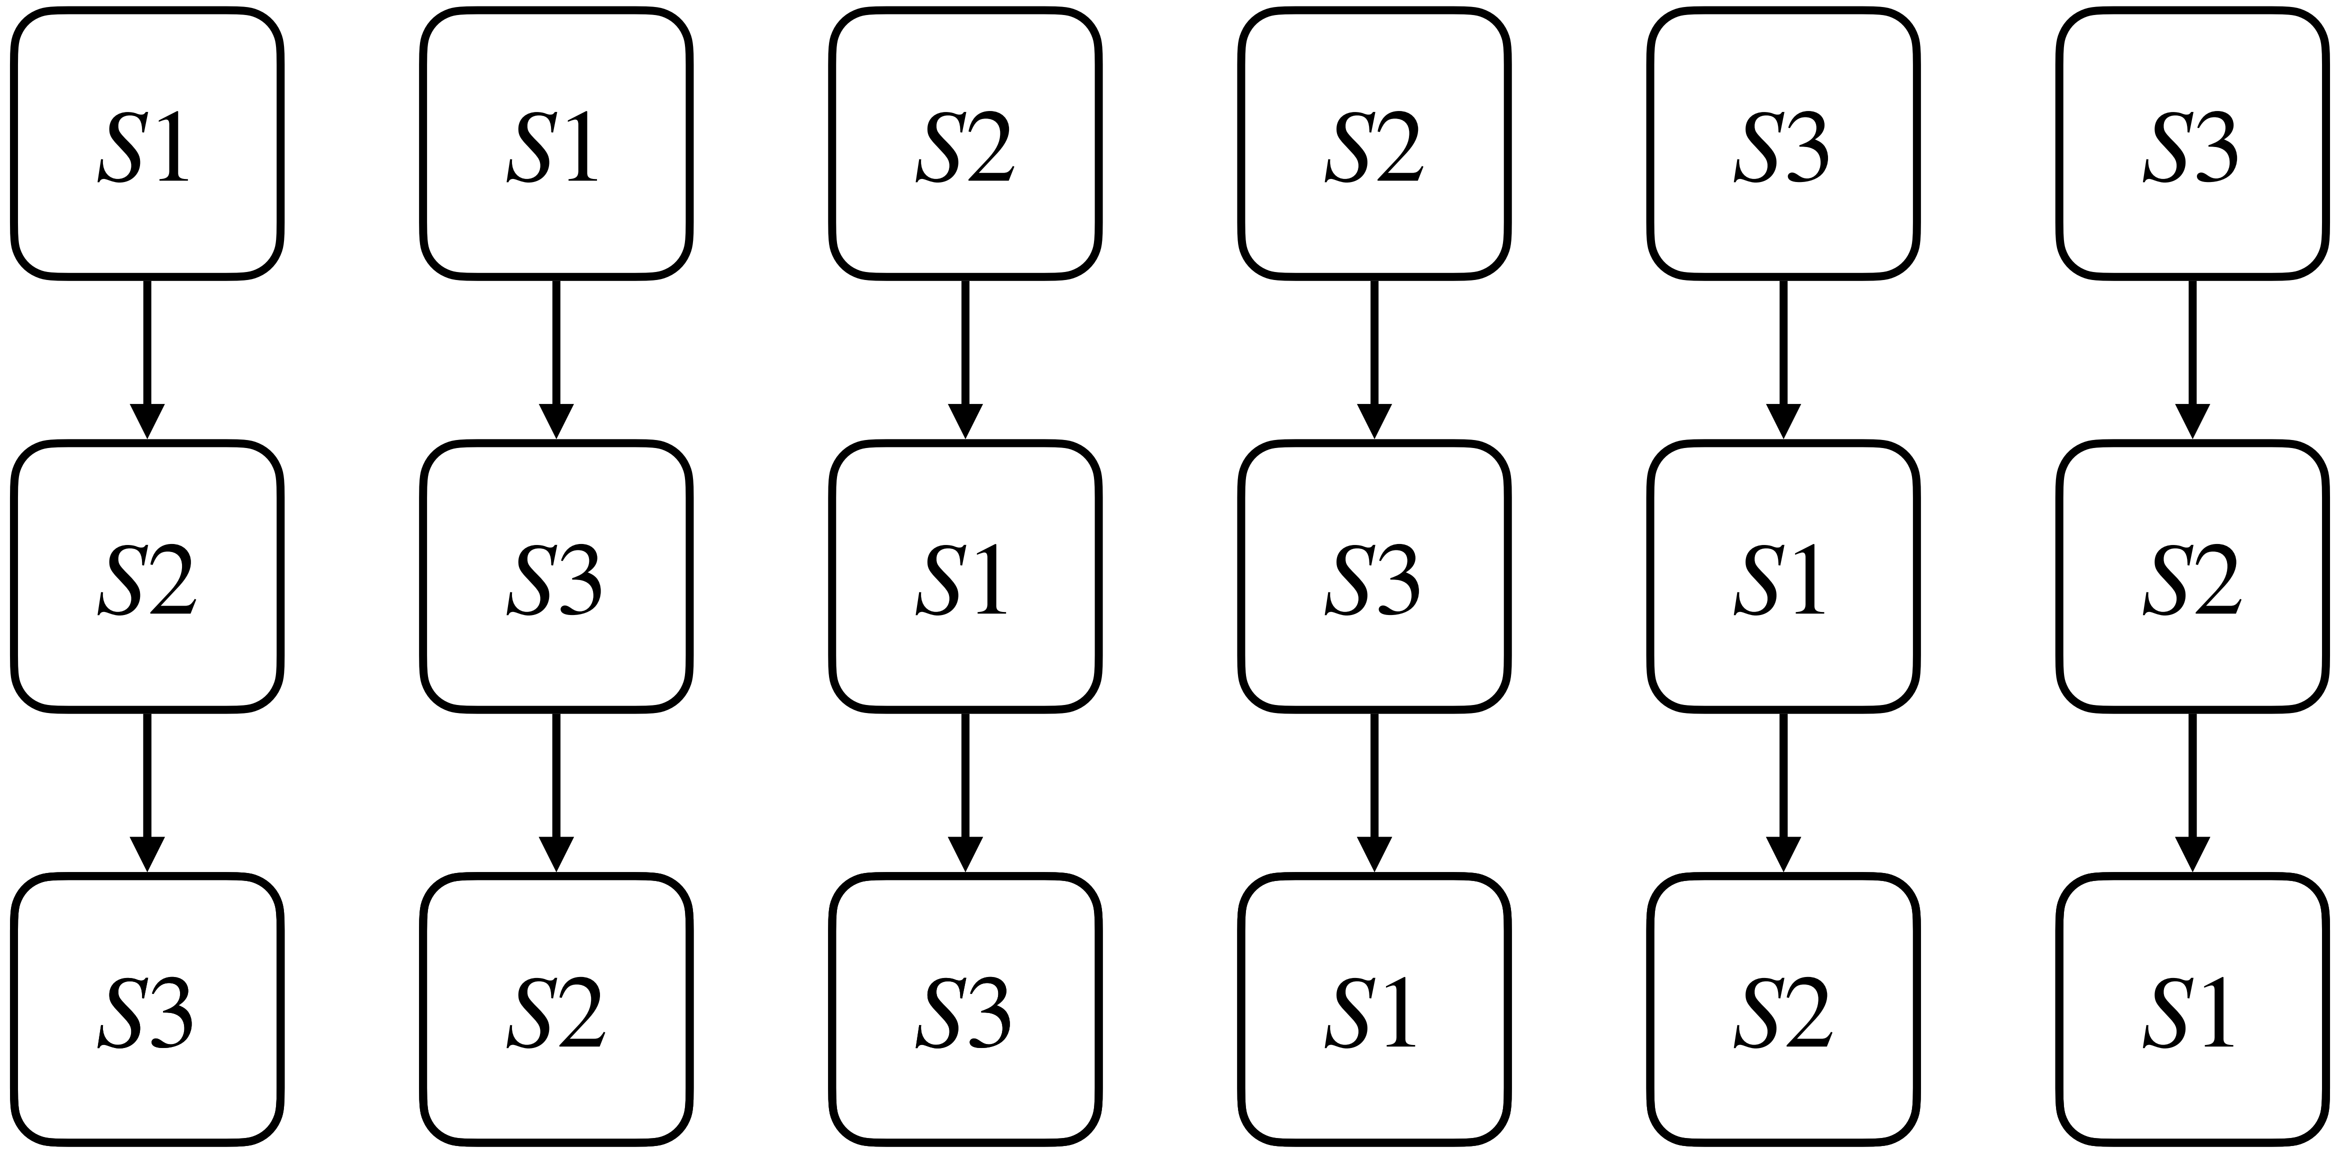

S^3T: Sequence Selection

Budget $B = 3$

L = 3, $L!$ = 6 sequences

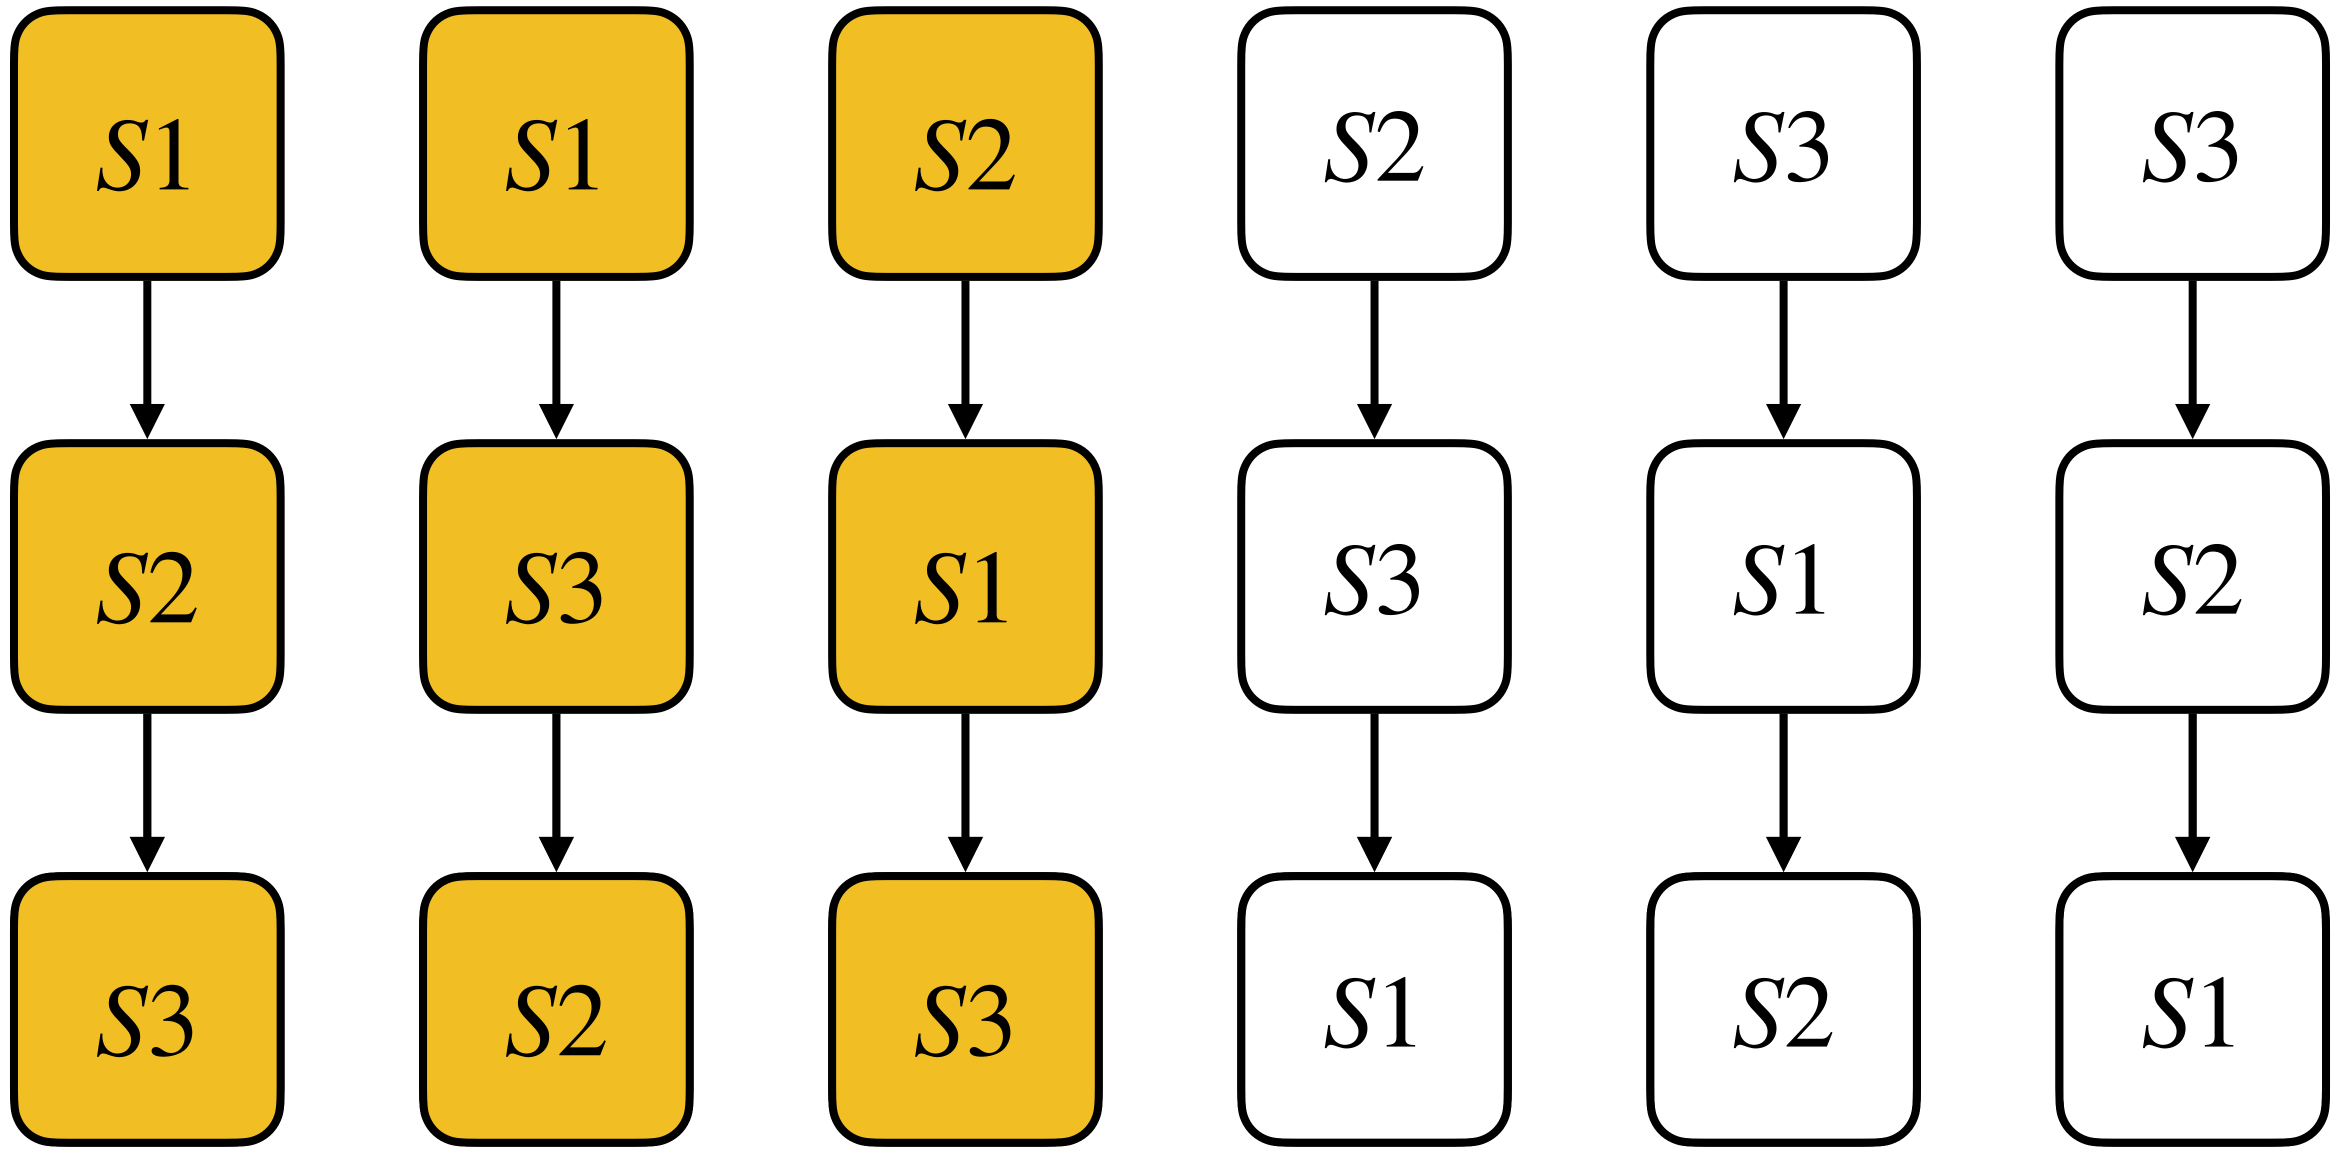

S^{3}T: Sequence Selection

Budget $B = 3$

$L = 3, L! = 6$ sequences

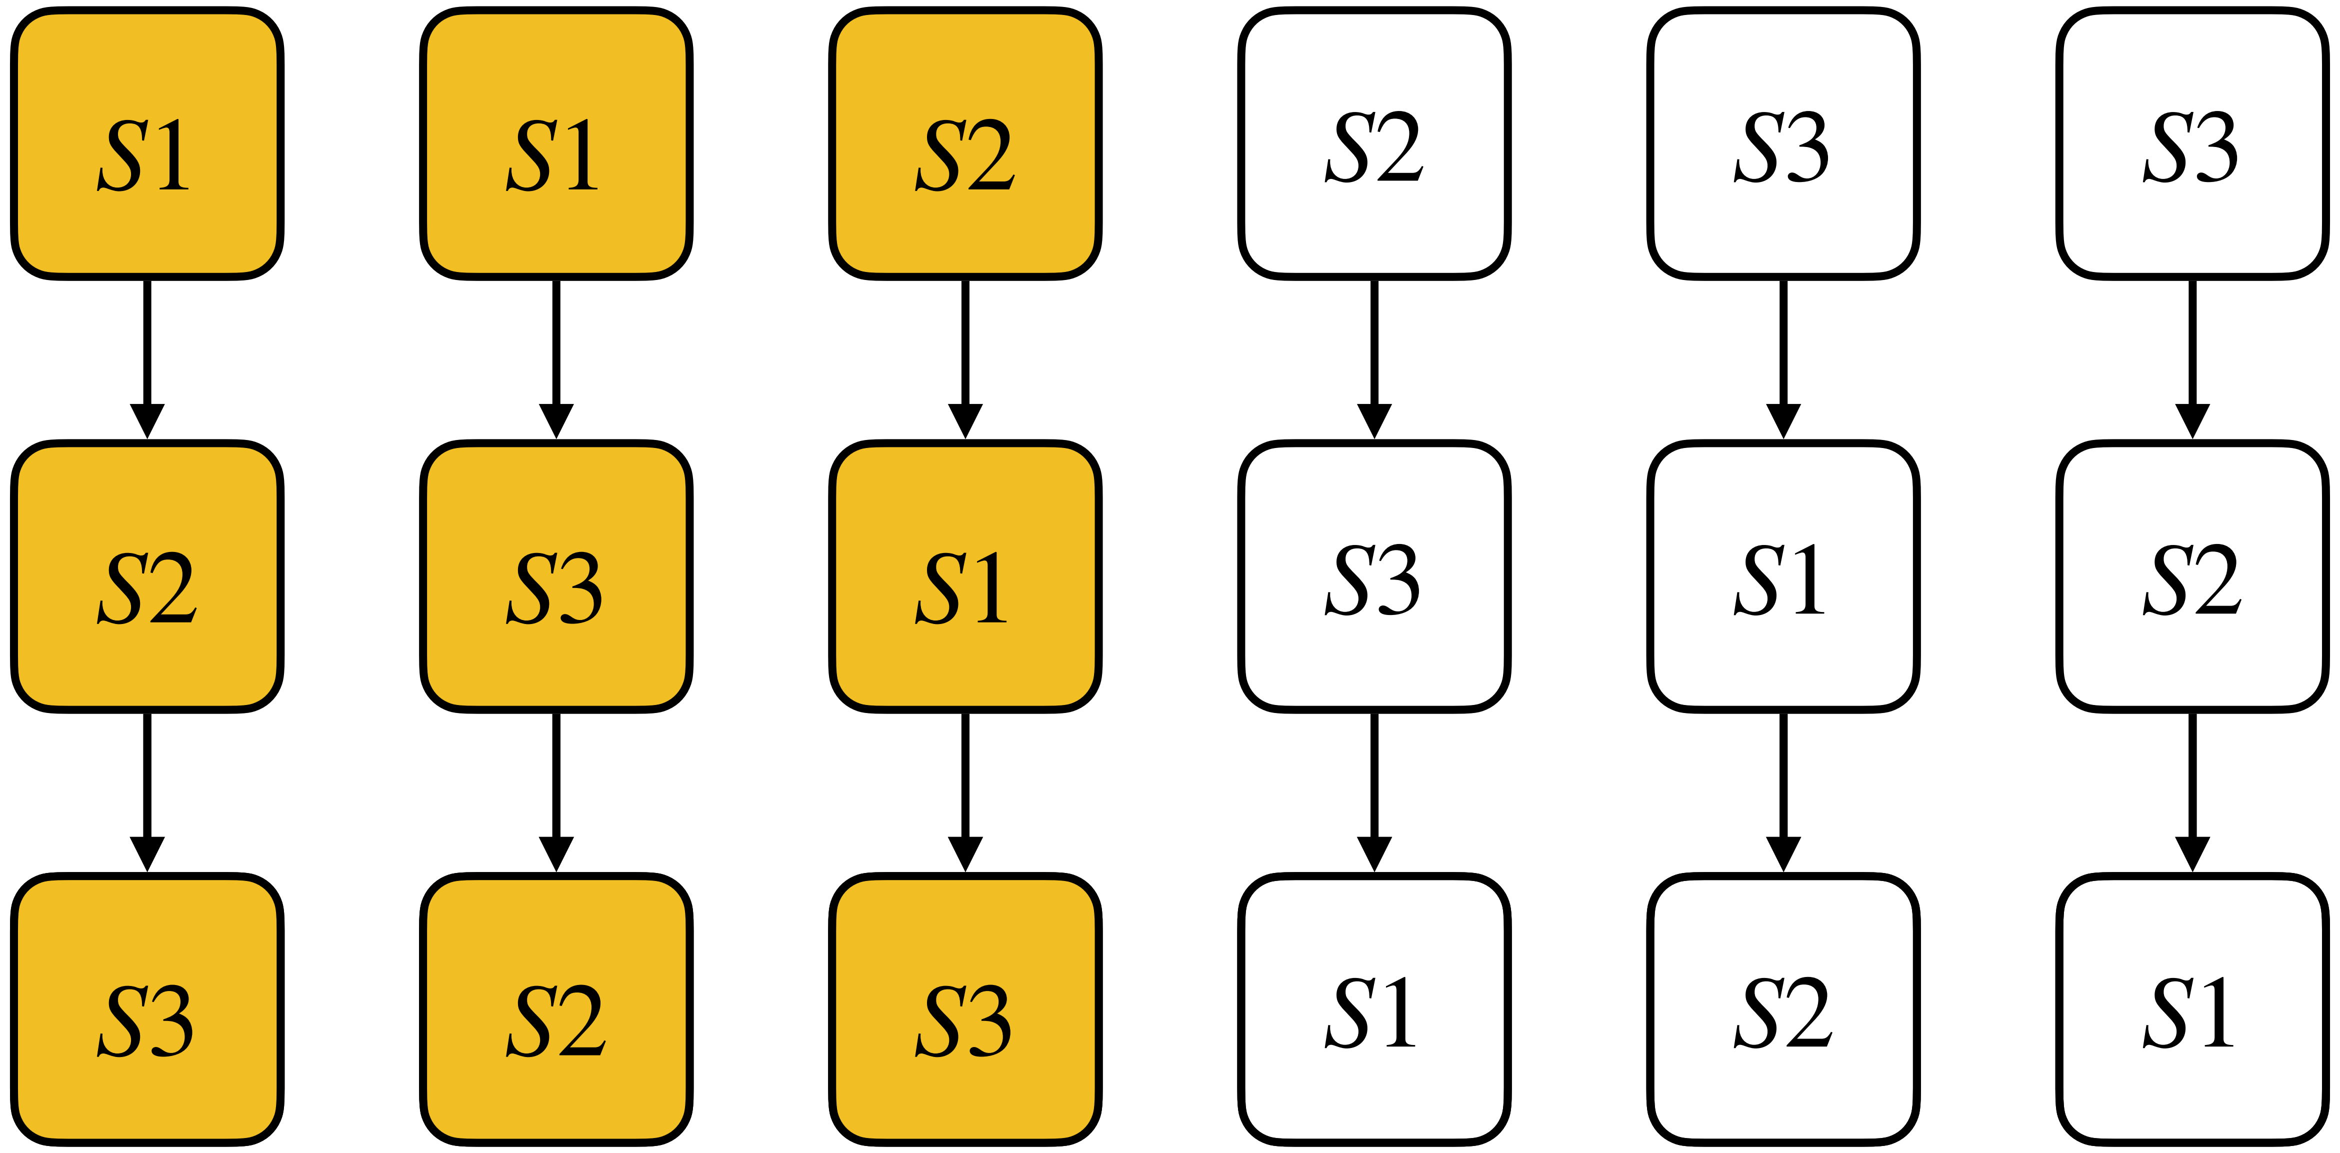

$S1$ $S2$ $S3$

$S1$ $S3$ $S2$

$S2$ $S1$ $S3$

$S2$ $S3$ $S1$

$S3$ $S1$ $S2$

$S3$ $S2$ $S1$

1. **Unknown Prior:** Iterative Cyclic Rotation 2. **Known Prior:** Bipartite Matching

S^{3}T: Sequence Selection

Budget $B = 3$

$L = 3, L! = 6$ sequences

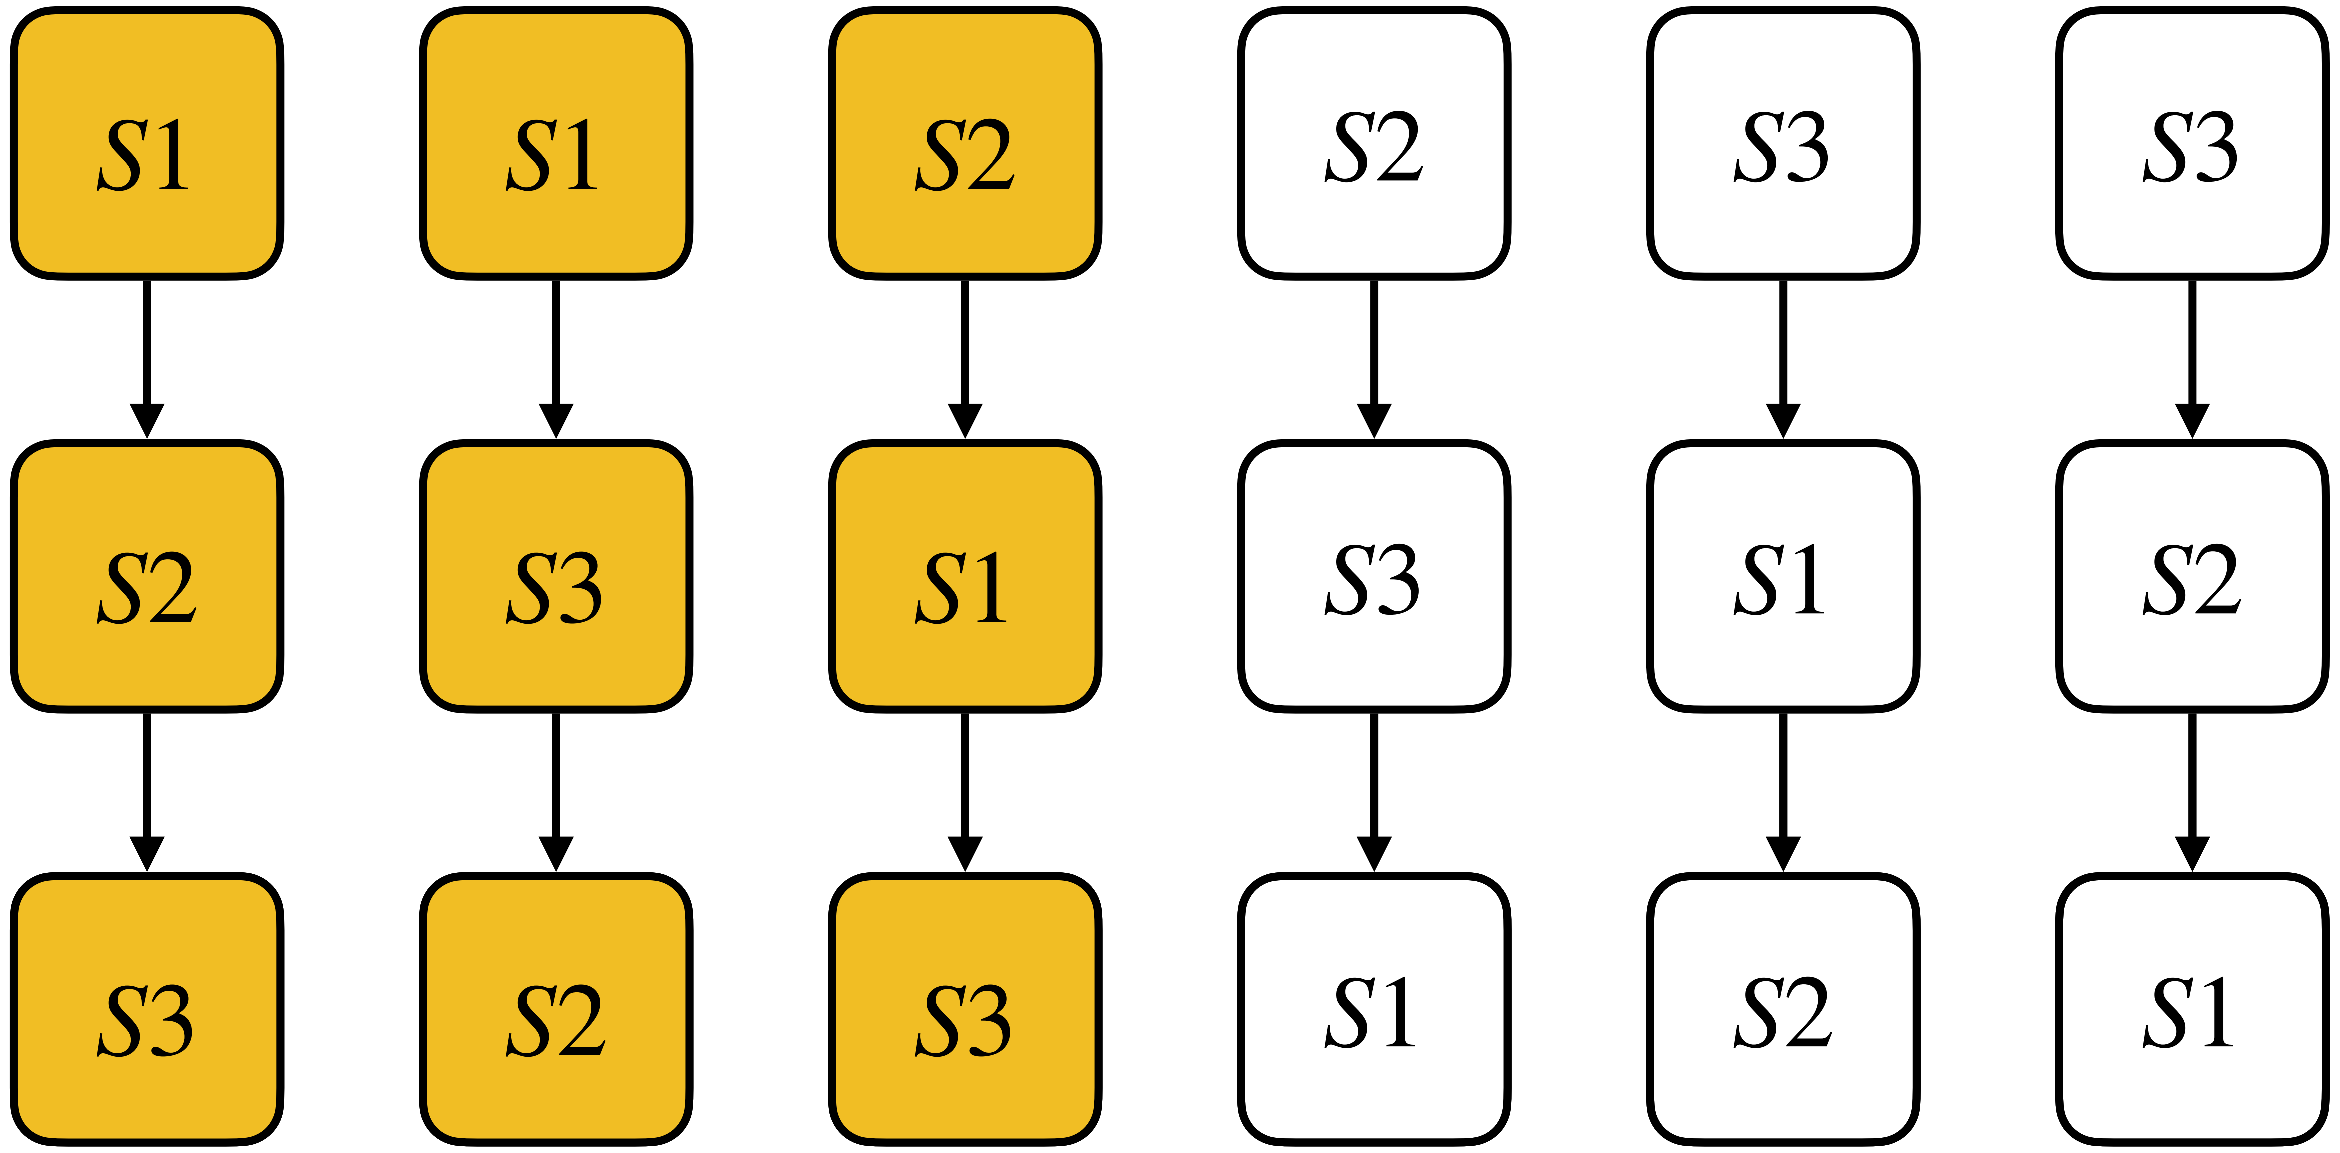

Leads to better **deletion guarantees**!

1. **Unknown Prior:** Iterative Cyclic Rotation 2. **Known Prior:** Bipartite Matching

S^{3}T Deletion Performance $(L = 4)$

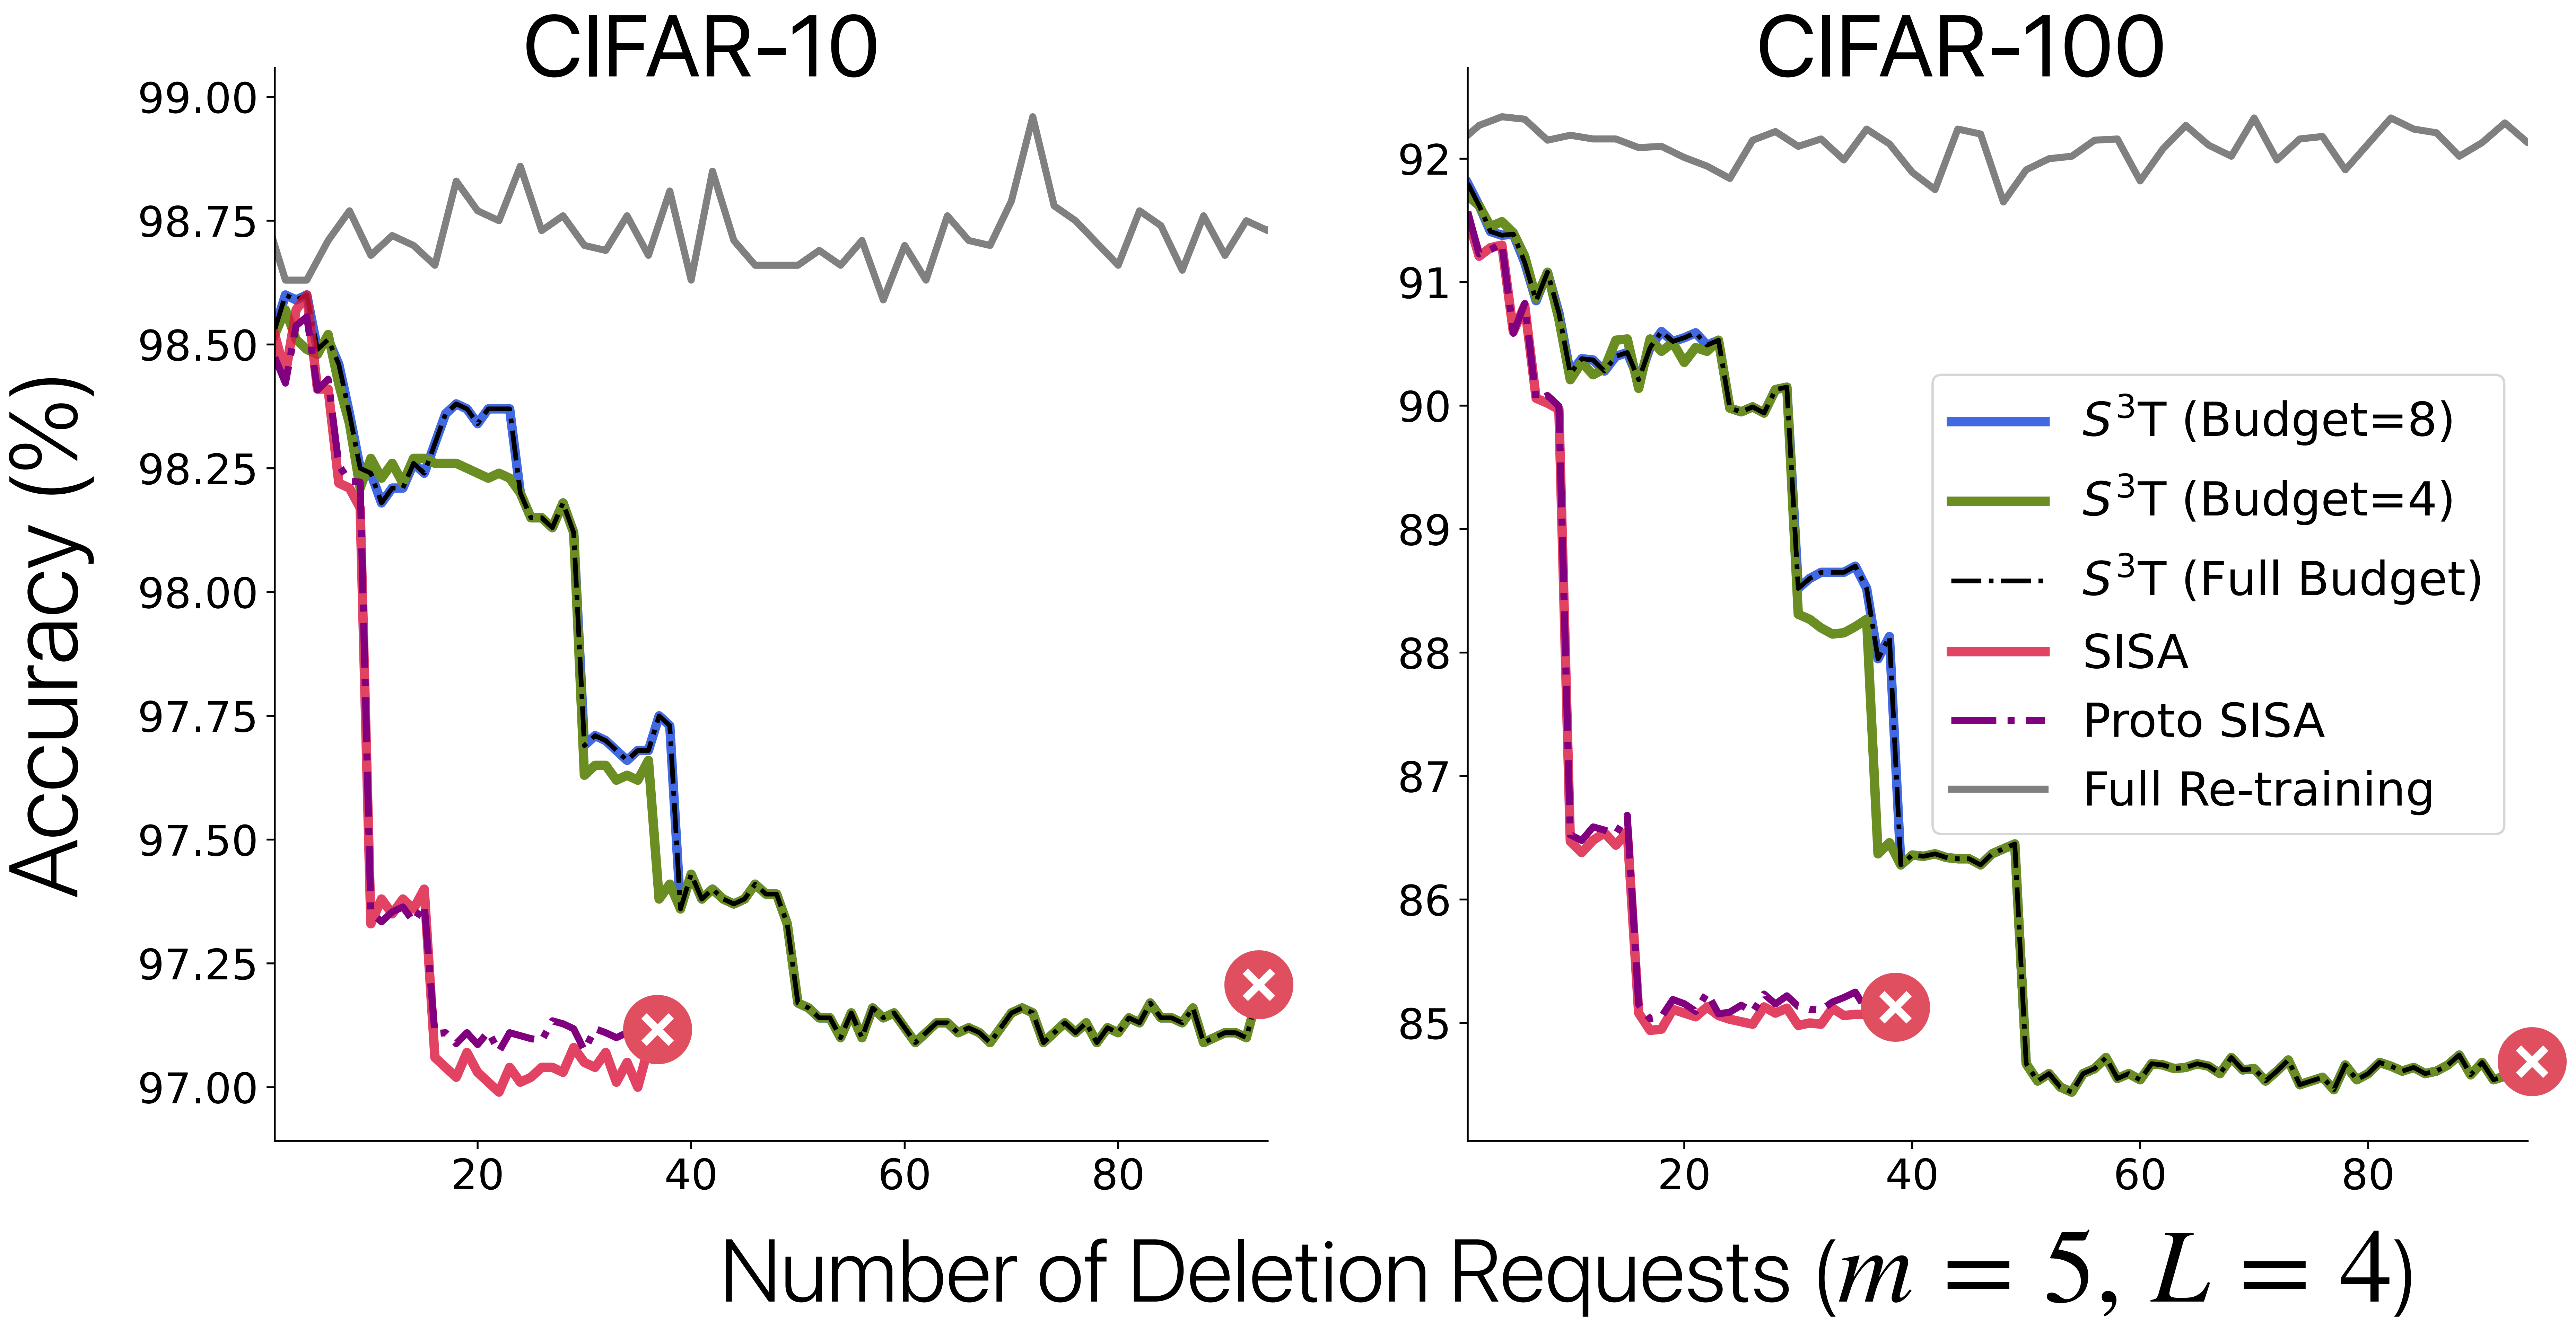

Summary

- We introduce an unlearning framework that achieves modularity using fine-tuning
- S^{3}T results in better theoretical guarantees about deletion requests
- In practice, S^{3}T can handle up to 4x more deletion requests than existing systems